# Data on Disk and B+-trees



## Storage Media: Players

Cache – fastest and most costly form of storage; volatile; managed by the computer system hardware.

#### Main memory:

- fast access (10s to 100s of nanoseconds; 1 nanosecond = 10<sup>-9</sup> seconds)
- generally too small (or too expensive) to store the entire database
- Volatile contents of main memory are usually lost if a power failure or system crash occurs.
- But... CPU operates only on data in main memory

# Storage Media: Players

#### Disk

- Primary medium for the long-term storage of data; typically stores entire database.
- random-access possible to read data on disk in any order, unlike magnetic tape
- Non-volatile: data survive a power failure or a system crash, disk failure less likely than them
- New technology: Solid State Disks and Flash disks

## Storage Media: Players

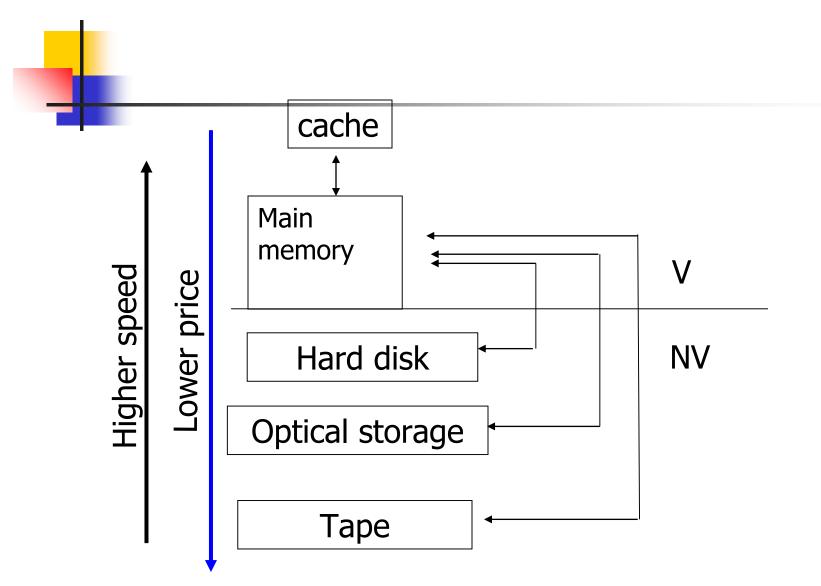
#### Optical storage

- non-volatile, data is read optically from a spinning disk using a laser
- CD-ROM (640 MB) and DVD (4.7 to 17 GB) most popular forms
- Write-one, read-many (WORM) optical disks used for archival storage (CD-R and DVD-R)
- Multiple write versions also available (CD-RW, DVD-RW, and DVD-RAM)
- Reads and writes are slower than with magnetic disk

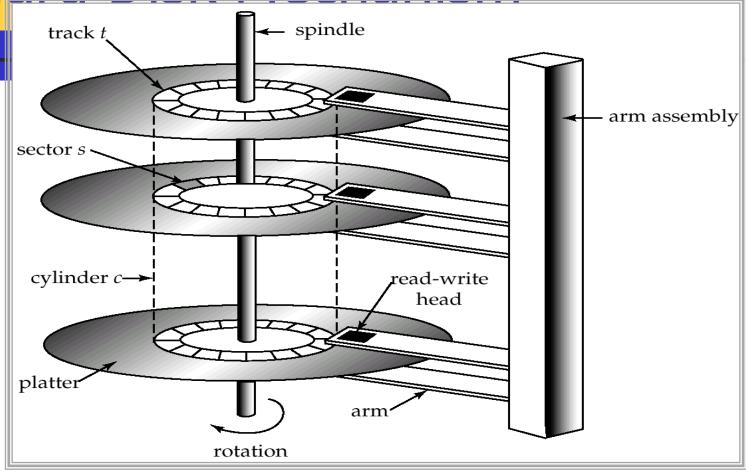
#### Tapes

- Sequential access (very slow)
- Cheap, high capacity

## Memory Hierarchy



Hard Disk Mechanism





#### Read-write head

- Positioned very close to the platter surface (almost touching it)
- Surface of platter divided into circular tracks
- Each track is divided into sectors.
  - A sector is the smallest unit of data that can be read or written.
- To read/write a sector
  - disk arm swings to position head on right track
  - platter spins continually; data is read/written as sector passes under head
- Block: a sequence of sectors
- Cylinder i consists of the track of all the platters

# 4

#### "Typical" Values

Diameter: 1 inch  $\rightarrow$  10 inches

Cylinders:  $100 \rightarrow 2000$ 

Surfaces: 1 or 2

(Tracks/cyl)  $2 \rightarrow 30$ 

Sector Size:  $512B \rightarrow 50K$ 

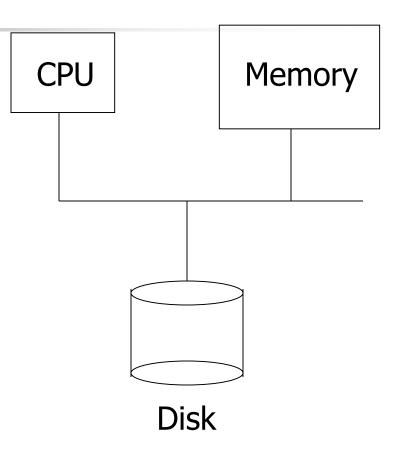
Capacity:  $4 \text{ MB} \rightarrow 8 \text{ TB}$ 

#### **Pata Access**

- Data Access Method (aka Index)
  - data structure that allows efficient access of data stored on hard disk for specific queries (query types)
- Query types:
  - Exact match
  - Range Query
  - more
- With New Data types and Applications =>
  - New Indexing and Query algorithms!!
  - New Systems!!!

### **Model of Computation**

- Data stored on disk(s)
- Minimum transfer unit: a page = b bytes or B records (or block)
- N records -> N/B = n pages
- I/O complexity: in number of pages



### I/O complexity

- ♣ An ideal index has space O(N/B), update overhead O(1) or O(log<sub>B</sub>(N/B)) and search complexity O( $\alpha$ /B) or O(log<sub>B</sub>(N/B) +  $\alpha$ /B) where  $\alpha$  is the number of records in the answer
  - But, sometimes CPU performance is also important... minimize cache misses -> don't waste CPU cycles





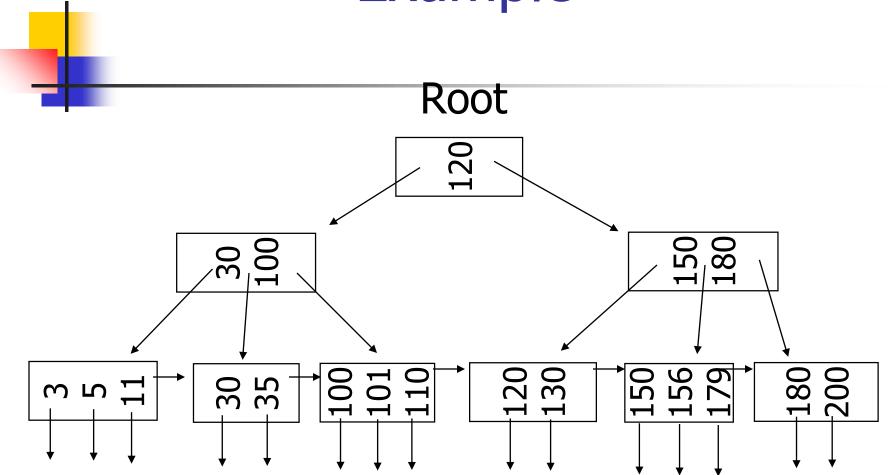
- Records must be <u>ordered</u> over an attribute, SSN, Name, etc.
- Queries: exact match and range queries over the indexed attribute: "find the name of the student with ID=087-34-7892" or "find all students with gpa between 3.00 and 3.5"

Optimal 1-dimensional Index for range queries!

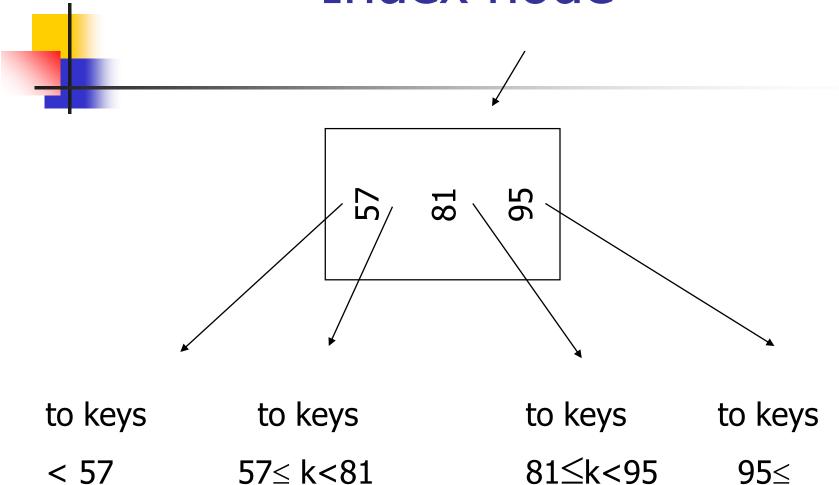
#### B+-tree:properties

- Insert/delete at log <sub>F</sub> (N/B) cost; keep tree height-balanced. (F = fanout)
- Minimum 50% occupancy (except for root). Each node contains  $\mathbf{d} <= \underline{m} <= 2\mathbf{d}$  entries/pointers.
- Two types of nodes: index nodes and data nodes; each node is 1 page (disk based method)

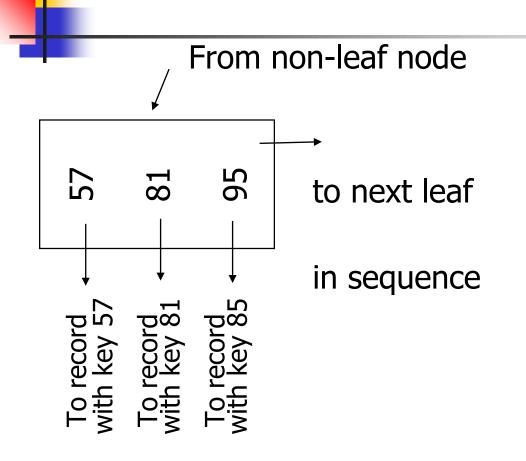
#### Example



#### Index node



#### Data node



```
Struct {
Key real;
Pointr *long;
} entry;
```

Node entry[B];

#### Insertion



- Find correct leaf L.
- Put data entry onto L.
  - If *L* has enough space, *done*!
  - Else, must *split L (into L and a new node L2)* 
    - Redistribute entries evenly, **copy up** middle key.
    - Insert index entry pointing to L2 into parent of L.
- This can happen recursively
  - To split index node, redistribute entries evenly, but **push up** middle key. (Contrast with leaf splits.)
- Splits "grow" tree; root split increases height.
  - Tree growth: gets <u>wider</u> or <u>one level taller at top</u>.

#### Deletion

- Start at root, find leaf L where entry belongs.
- Remove the entry.
  - If L is at least half-full, done!
  - If L has only d-1 entries,
    - Try to re-distribute, borrowing from sibling (adjacent node with same parent as L).
    - If re-distribution fails, merge L and sibling.
- If merge occurred, must delete entry (pointing to L or sibling) from parent of L.
- Merge could propagate to root, decreasing height.

# 4

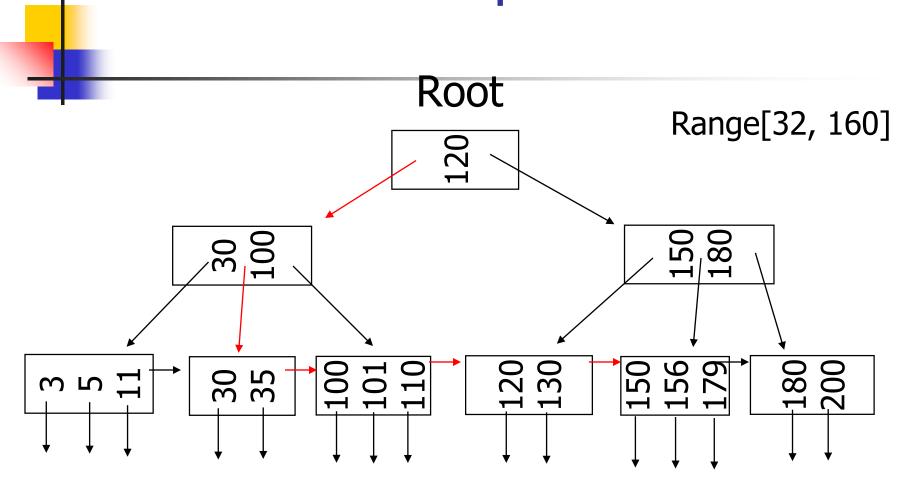
#### Characteristics

Optimal method for 1-d range queries:

Space: O(N/B), Updates: O( $log_B(N/B)$ ), Query:O( $log_B(N/B) + \alpha/B$ )

- Space utilization: 67% for random input
- B-tree variation: index nodes store also pointers to records

## Example



#### Other issues



- Internal node architecture [Lomet01]:
  - Reduce the overhead of tree traversal.
    - Prefix compression: In index nodes store only the prefix that differentiate consecutive sub-trees. Fanout is increased.
- Cache sensitive B+-tree
  - Place keys in a way that reduces the cache faults during the binary search in each node.
  - Eliminate pointers so a cache line contains more keys for comparison.



# **External Memory Hashing**



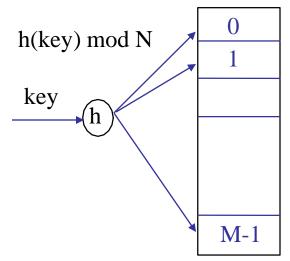
## Hashing

- Hash-based indices are best for exact match queries. Faster than B+-tree!
- Typically 1-2 I/Os per query where a B+tree requires 4-5 I/Os
- But, cannot answer range queries...

#### Idea



- Use a function to direct a record to a page
- h(k) mod M = bucket to which data entry with key k belongs. (M = # of buckets)

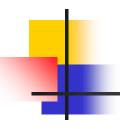


Primary bucket pages

#### Design decisions: Functions

- h(x) = ((a\*x+b) mod p) mod M
 a>0, a,b <p, p is a prime, and p>M
 (this is a Universal Hash family for different (a,b) values)

- $h(x) = [fractional-part-of(x*\varphi)]*M$ ,  $\varphi$ : golden ratio (0.618... = (sqrt(5)-1)/2)
- Size of hash table M
- Overflow handling: open addressing or chaining: problem in dynamic databases



## Dynamic hashing schemes

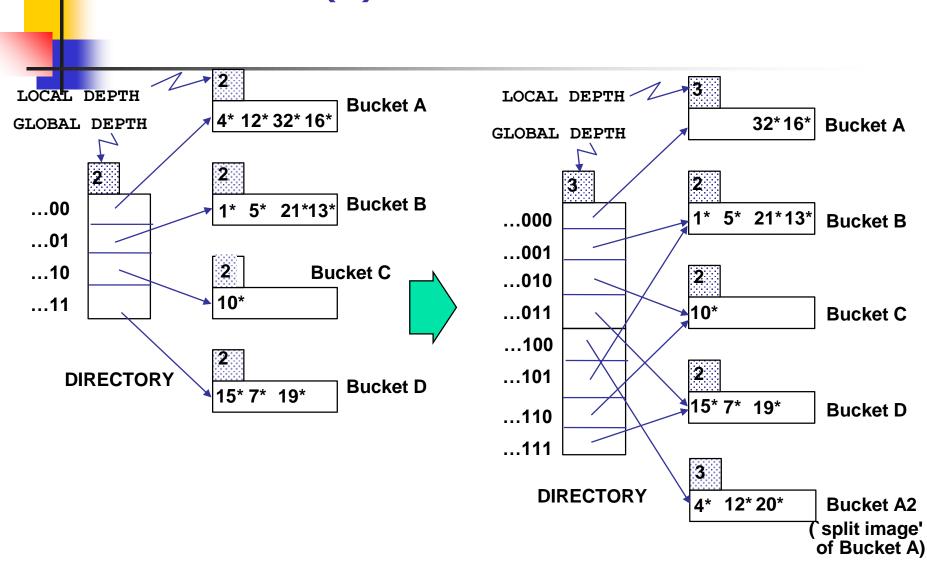
- <u>Extensible hashing</u>: uses a directory that grows or shrinks depending on the data distribution. No overflow buckets
- <u>Linear hashing</u>: No directory. Splits buckets in linear order, uses overflow buckets



## **Extensible Hashing**

- Bucket (primary page) becomes full. Why not reorganize file by doubling # of buckets (changing the hash function)?
  - Reading and writing all pages is expensive!
  - Idea: Use <u>directory of pointers to buckets</u>, double # of buckets by doubling the directory, splitting just the bucket that overflowed!
  - Directory much smaller than file, so doubling it is much cheaper. Only one page of data entries is split.
  - Trick lies in how hash function is adjusted!

#### Insert $h(k) = 20\ 10100 \rightarrow 00$



### Linear Hashing

- This is another dynamic hashing scheme, alternative to Extensible Hashing.
  - Motivation: Ext. Hashing uses a directory that grows by doubling... Can we do better? (smoother growth)
  - LH: split buckets from left to right, regardless of which one overflowed (simple, but it works!!)

# Linear Hashing (Contd.)

- Directory avoided in LH by using overflow pages.
   (chaining approach)
  - Splitting proceeds in `rounds'. Round ends when all M<sub>R</sub> initial (for round R) buckets are split. Buckets 0 to Next-1 have been split; Next to M<sub>R</sub> yet to be split.
  - Current round number is Level.
  - Search: To find bucket for data entry r, find h<sub>Level</sub>(r):
    - If  $\mathbf{h}_{Level}(r)$  in range `Next to  $M_R$ ', r belongs here.
    - Else, r could belong to bucket  $\mathbf{h}_{Level}(r)$  or bucket  $\mathbf{h}_{Level}(r) + \mathbf{M}_{R}$ ; must apply  $\mathbf{h}_{Level+1}(r)$  to find out.

# 4

## Linear Hashing: Example

```
Initially: h(x) = x \mod M (M=4 here)
Assume 3 records/bucket
```

Insert 17 = 17 mod 4 
$$\longrightarrow$$
 1

**Bucket id** 

$$hi(x) = x \mod 2^{Level} * M$$

Level=0

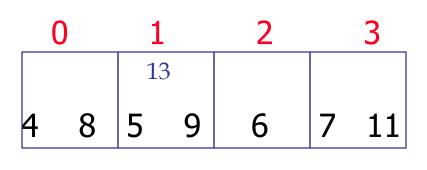
# 4

### Linear Hashing: Example

Initially:  $h(x) = x \mod N \pmod{N=4 \text{ here}}$ Assume 3 records/bucket Overflow for Bucket 1

Insert 17 = 17 mod 4 
$$\longrightarrow$$
 1

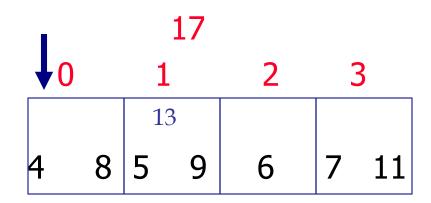
**Bucket** id



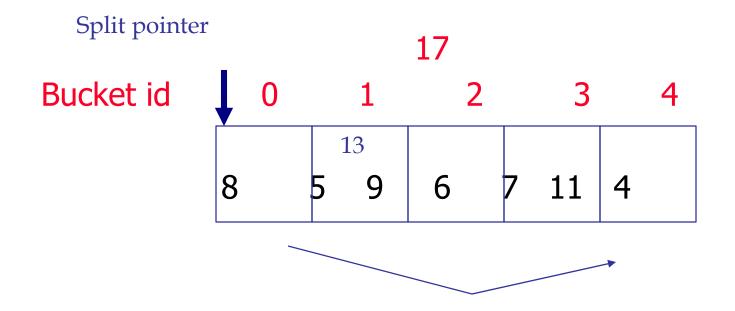
Split bucket 0, anyway!!

To split bucket 0, use another function h1(x):  $h0(x) = x \mod N$ ,  $h1(x) = x \mod (2*N)$ 


Split pointer



To split bucket 0, use another function h1(x):  $h0(x) = x \mod N$ ,  $h1(x) = x \mod (2*N)$ 




To split bucket 0, use another function h1(x):  $h0(x) = x \mod N$ ,  $h1(x) = x \mod (2*N)$ 



# 1


## Linear Hashing: Example

$$h0(x) = x \mod N$$
,  $h1(x) = x \mod (2*N)$ 




$$h0(x) = x \mod N$$
,  $h1(x) = x \mod (2*N)$ 

Bucket id



#### Linear Hashing: Search

 $h0(x) = x \mod N$  (for the un-split buckets)  $h1(x) = x \mod (2*N)$  (for the split ones)



#### Linear Hashing: Search

 $h1(x) = x \mod 8$  (for the un-split buckets)  $h2(x) = x \mod 16$  (for the split ones)

0	1	2	3	4	5	6	7
	17						
8	9	6	3 11	4	13 5		15 7

After we split the Nth bucket (3), we reset the Next pointer to 0 and we start a new round. The two hash functions are now h1 and h2.

Level =1

#### Linear Hashing: Search

Algorithm for Search:

Search(k)

- 1 b = h0(k)
- 2 if b < split-pointer then
- b = h1(k)
- 4 read bucket b and search there

# References

[Litwin80] Witold Litwin: Linear Hashing: A New Tool for File and Table Addressing. VLDB 1980: 212-223

http://www.cs.bu.edu/faculty/gkollios/ada01/Papers/linear-hashing.PDF