



Spatial Indexing I


Point Access Methods

Many slides are based on slides provided
by Prof. Christos Faloutsos (CMU) and
Prof. Evimaria Terzi (BU)



The problem: Range query (or range reporting)

- Given a point set and a rectangular query, find the points enclosed in the query
- We allow insertions/deletions on line





Outline




- Grid file
- kd-tree based methods
- Space filing curves
 - Quadtrees



Grid File

- Hashing methods for multidimensional points (extension of Extensible hashing)
- Idea: Use a grid to partition the space → each cell is associated with one page
- Two disk access principle (exact match)


Grid File



- Start with one bucket for the whole space.
- Select dividers along each dimension.
Partition space into cells
- Dividers cut all the way.



Grid File




- Each cell corresponds to 1 disk page.
- Many cells can point to the same page.
- Cell directory potentially exponential in the number of dimensions



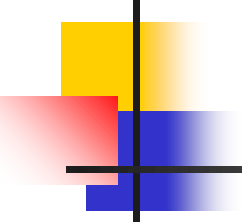
Grid File Implementation

- Dynamic structure using a grid directory
 - Grid array: a 2 dimensional array with pointers to buckets (this array can be large, disk resident) $G(0, \dots, nx-1, 0, \dots, ny-1)$
 - Linear scales: Two 1 dimensional arrays that used to access the grid array (main memory) $X(0, \dots, nx-1), Y(0, \dots, ny-1)$

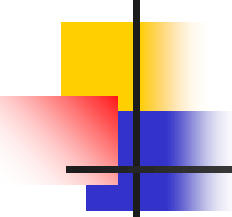
Example



Grid File Search

- 
- Exact Match Search: at most 2 I/Os assuming linear scales fit in memory.
 - First use linear scales to determine the index into the cell directory
 - access the cell directory to retrieve the bucket address (may cause 1 I/O if cell directory does not fit in memory)
 - access the appropriate bucket (1 I/O)
 - Range Queries:
 - use linear scales to determine the index(es) into the cell directory.
 - Access the cell directory to retrieve the bucket addresses of buckets to visit.
 - Access the buckets.

Grid File Insertions

- 
- Determine the bucket into which insertion must occur.
 - If space in bucket, insert.
 - Else, split bucket
 - how to choose a good dimension to split?
 - If bucket split causes a cell directory to split do so and adjust linear scales.
 - insertion of these new entries potentially requires a complete reorganization of the cell directory---expensive!!!



Grid File Deletions

- Deletions may decrease the space utilization.
Merge buckets
- We need to decide which cells to merge and a merging threshold
- Buddy system and neighbor system
 - A bucket can merge with only one buddy in each dimension
 - Merge adjacent regions if the result is a rectangle



Tree-based PAMs

- Most of tb-PAMs are based on kd-trees
- kd-tree (J. Bentley, 1975) is a main memory binary tree for indexing k-dimensional points
 - Needs to be adapted for the disk model
- Levels rotate among the dimensions, partitioning the space based on a value for that dimension
- kd-tree is not necessarily balanced for dynamic datasets



2-dimensional kd-trees

- A data structure to support range queries in $\mathbf{R^2}$
 - Not the most efficient solution in theory
 - Used a lot in practice
- Preprocessing time: **$O(n \log n)$**
- Space complexity: **$O(n)$**
- Query time: **$O(n^{1/2} + P)$ (worst case)**
 - **P** number of points in the answer



2-dimensional kd-trees

- Idea: Split the point set alternating by x-coordinate and by y-coordinates
 - split by x-coordinate: split by a vertical line that has half the points left or on, and half right
 - split by y-coordinate: split by a horizontal line that has half the points below or on, and half above
- We get a binary tree:
 - Size **$O(n)$**
 - Depth **$O(\log n)$**
 - Construction time **$O(n \log n)$**



BuildKdtree

Algorithm BuildKdTree($P, depth$)

if P contains only one point

 then return a leaf storing this point

else if $depth$ is even

 then Split P with a vertical line l through the median x -coordinate

 into P_1 (left of or on l) and P_2 (right of l)


else Split P with a horizontal line l through the median y -coordinate

 into P_1 (below or on l) and P_2 (above l)

 BuildKdTree($P_1, depth + 1$)




Construction of kd-trees







Construction of kd-trees




Construction of kd-trees





Construction of kd-trees



Construction of kd-trees



The complete kd-tree





Construction cost


- Finding median is $O(n)$
- Also: $T(1) = O(1)$

$$T(n) = 2 T(n/2) + O(n)$$

(Master Theorem, MT)

$$\Rightarrow T(n) = O(n \log n)$$

Region of node **v**



Region(v) : the subtree rooted at **v** stores the points in black dots



Searching in kd-trees

- Range-searching in **2-d**
 - Given a set of **n** points, build a data structure that for any query rectangle **R** reports all point in **R**




kd-tree: range queries

- Recursive procedure starting from **$v = \text{root}$**
- **Search** (v, R)
 - If v is a leaf, then report the point stored in v if it lies in R
 - Otherwise, if **Reg**(v) is contained in R , report all points in the **subtree**(v)
 - Otherwise:
 - If **Reg**(**left**(v)) intersects R , then **Search**(**left**(v), R)
 - If **Reg**(**right**(v)) intersects R , then **Search**(**right**(v), R)


Query time analysis



- We will show that **Search** takes at most $O(n^{1/2} + P)$ time, where P is the number of reported points
 - The total time needed to report all points in all sub-trees is $O(P)$
 - We just need to bound the number of nodes v such that **region**(v) intersects R but is not contained in R (i.e., boundary of R intersects the boundary of **region**(v))
 - **gross overestimation**: bound the number of **region**(v) which are crossed by any of the 4 horizontal/vertical lines



- **$Q(n)$** : max number of regions in an n -point kd-tree intersecting a (say, vertical) line?



- If ℓ intersects **region(v)** (due to vertical line splitting), then after two levels it intersects **2** regions (due to 2 vertical splitting lines)
- The number of regions intersecting ℓ is **$Q(n)=2+2Q(n/4)$** , **MT**
 \rightarrow **$Q(n)=(n^{1/2})$**



Query time (Cont'd)

- Range-searching in **2d kd-tree**
- A range has four sides. In the worst case, every side intersects $O(n^{1/2})$ regions (nodes to follow)
- So, total query time is:
$$Q(n) = (n^{1/2} + p)$$



d-dimensional kd-trees

- A data structure to support range queries in \mathbf{R}^d
- Preprocessing time: **$O(n \log n)$**
- Space complexity: **$O(n)$**
- Query time: **$O(n^{1-1/d} + P)$**



Construction of the **d**-dimensional kd-trees

- The construction algorithm is similar as in **2-d**
- At the root we split the set of points into two subsets of same size by a hyper-plane vertical to \mathbf{x}_1 -axis
- At the children of the root, the partition is based on the second coordinate: \mathbf{x}_2 -coordinate
- At depth **d**, we start all over again by partitioning on the first coordinate
- The recursion stops until there is only one point left, which is stored as a leaf

External memory kd-trees (kdB-tree)



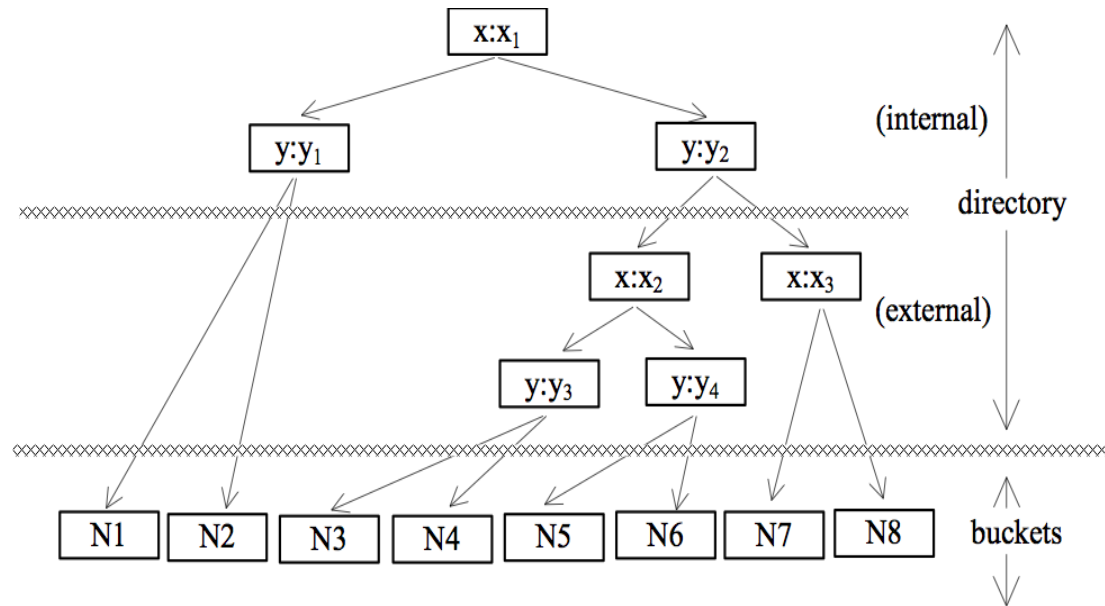
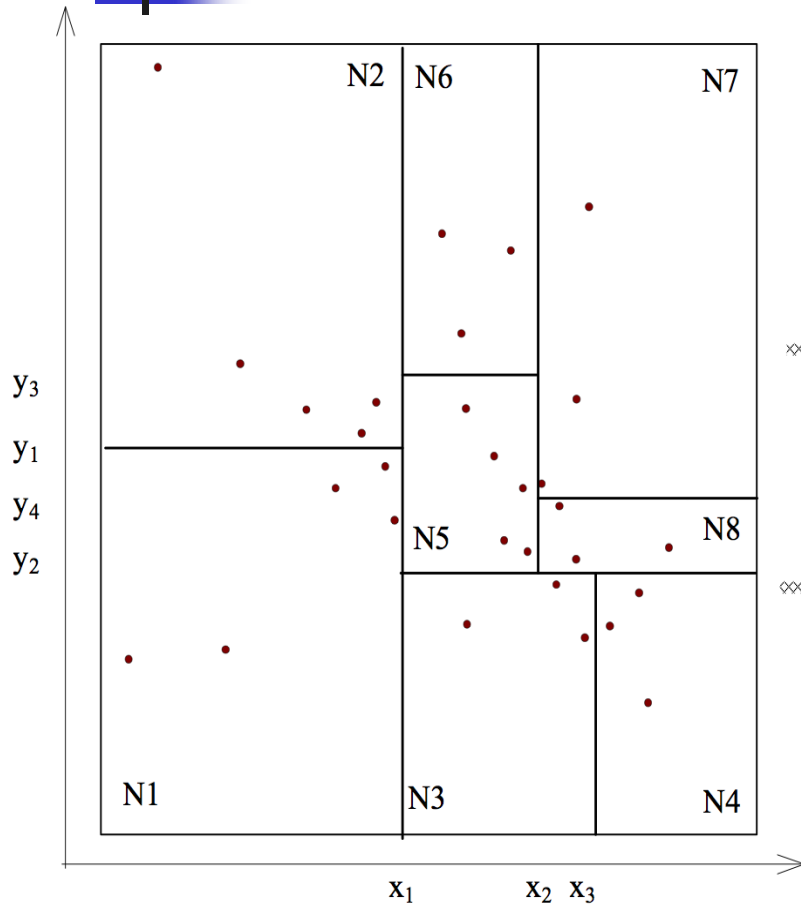
- Pack many interior nodes (forming a subtree) into a block using BFS-traversal.
 - it may not be feasible to group nodes at lower level into a block productively.
 - Many interesting papers on how to optimally pack nodes into blocks recently published.
- Similar to B-tree, tree nodes split many ways instead of two ways
 - insertion becomes quite complex and expensive.
 - No storage utilization guarantee since when a higher level node splits, the split has to be propagated all the way to leaf level resulting in many empty blocks.



LSD-tree

- Local Split Decision – tree
- Use kd-tree to partition the space. Each partition contains up to B points. The kd-tree is stored in main-memory.
- If the kd-tree (directory) is large, we store a sub-tree on disk
- Goal: the structure must remain balanced: external balancing property

Example: LSD-tree





LSD-tree: main points

- Split strategies:
 - Data dependent
 - Distribution dependent
- Paging algorithm
- Two types of splits: bucket splits and internal node splits



PAMs

- Point Access Methods
 - Multidimensional Hashing: Grid File
 - Exponential growth of the directory
 - Hierarchical methods: kd-tree based
 - Storing in external memory is tricky but possible
 - Space Filling Curves: Z-ordering
 - Map points from 2-dimensions to 1-dimension.
Use a B+-tree to index the 1-dimensional points




Z-ordering

- Basic assumption: Finite precision in the representation of each co-ordinate, K bits (2^K values)
- The address space is a square (image) and represented as a $2^K \times 2^K$ array
- Each element is called a pixel

Z-ordering

- Impose a linear ordering on the pixels of the image \rightarrow 1 dimensional problem



$$Z_A = \text{shuffle}(x_A, y_A) = \text{shuffle}("01", "11")$$

$$= 0111 = (7)_{10}$$

$$Z_B = \text{shuffle}("01", "01") = 0011$$




Z-ordering

- Given a point (x, y) and the precision K find the pixel for the point and then compute the z-value
- Given a set of points, use a B+-tree to index the z-values
- A range (rectangular) query in 2-d is mapped to a set of ranges in 1-d

Queries

- Find the z-values that contained in the query and then the ranges



$Q_A \rightarrow$ range [4, 7]

$Q_B \rightarrow$ ranges [2,3] and [8,9]




Hilbert Curve

- We want points that are close in 2d to be close in the 1d
- Note that in 2d there are 4 neighbors for each point where in 1d only 2.
- Z-curve has some “jumps” that we would like to avoid
- Hilbert curve avoids the jumps :
recursive definition




Hilbert Curve- example


- It has been shown that in general Hilbert is better than the other space filling curves for retrieval [Jag90]
- H_i (order- i) Hilbert curve for $2^i \times 2^i$ array



H1



H2



... $H(n+1)$

Handling Regions


- A region breaks into one or more pieces, each one with different z-value
- Works for raster representations (pixels)
- We try to minimize the number of pieces in the representation: precision/space overhead trade-off

$$Z_{R1} = 0010 = (2)$$

$$Z_{R2} = 1000 = (8)$$

$$Z_G = 11$$


(“11” is the common prefix)






Z-ordering for Regions

- Break the space into 4 equal quadrants: level-1 blocks
- Level-i block: one of the four equal quadrants of a level-(i-1) block
- Pixel: level-K blocks, image level-0 block
- For a level-i block: all its pixels have the same prefix up to $2i$ bits; the z-value of the block




Quadtree

- Object is recursively divided into blocks until:
 - Blocks are homogeneous
 - Pixel level
- Quadtree: '0' stands for S and W
'1' stands for N and E



11
1001
1011






Region Quadrees

- Implementations
 - FL (Fixed Length)
 - FD (Fixed length-Depth)
 - VL (Variable length)
- Use a B+-tree to index the z-values and answer range queries



Linear Quadtree (LQ)

- Assume we use n -bits in each dimension (x, y) (so we have $2^n \times 2^n$ pixels)
- For each object O , compute the z -values of this object: $z_1, z_2, z_3, \dots, z_k$ (each value can have between 0 and $2n$ bits)
- For each value z_i we append at the end the level \perp of this value (level $\perp = \log(|z_i|)$)
- We create a value with $2n+1$ bits for each z -value and we insert it into a B+-tree ($l = \log_2(h)$)



Z-value, 1 | Morton block

A: 00, 01 = 00000001


B: 0110, 10 = 01100010

C: 111000, 11 = 11100011

n=3

A:1, B:98, C: 227

Insert into B+-tree using Mb





Query Alg

WindowQ(query w, quadtree block b)

{ Mb = Morton block of b;

If b is totally enclosed in w {

 Compute Mbmax

 Use B+-tree to find all objects with M values between $M_b \leq M \leq M_{bmax}$

 add to result

} else {

 Find all objects with Mb in the B+-tree

 add to result

 Decompose b into four quadrants sw, nw, se, ne

 For child in {sw, nw, se, ne}

 if child overlaps with w

 WindowQ(w, child)

 }



}



z-ordering - analysis

Q: How many pieces ('quad-tree blocks') per region?

A: proportional to perimeter (surface etc)








z-ordering - analysis

(How long is the coastline, say, of Britain?

Paradox: The answer changes with the yardstick -> fractals ...)







Unit: 200 km, 100 km and 50 km in length.

The resulting coastline is about 2350 km, 2775 km and 3425 km



z-ordering - analysis

Q: Should we decompose a region to full detail (and store in B-tree)?







z-ordering - analysis

Q: Should we decompose a region to full detail (and store in B-tree)?



A: NO! approximation with 1-5 pieces/z-values is best [Orenstein90]





z-ordering - analysis

Q: how to measure the 'goodness' of a curve?







z-ordering - analysis

Q: how to measure the 'goodness' of a curve?

A: e.g., avg. # of runs, for range queries



4 runs



3 runs

(#runs ~ #disk accesses on B-tree)



z-ordering - analysis

Q: So, is Hilbert really better?

A: 27% fewer runs, for 2-d (similar for 3-d)

Q: are there formulas for #runs, #of
quadtree blocks etc?

A: Yes, see a paper by [Jagadish '90]

H.V. Jagadish. Linear clustering of objects with multiple attributes.

SIGMOD 1990. <http://www.cs.ucr.edu/~tsotras/cs236/W15/hilbert-curve.pdf>