




Spatial Indexing I

R-trees



Problem

- Given a collection of geometric objects (points, lines, polygons, ...)
- organize them on disk, to answer efficiently spatial queries (range, nn, etc)



R-tree

- In multidimensional space, there is no unique ordering!
Not possible to use B+-trees☹
- [Guttman 84] R-tree!
- Group objects close in space in the same node
 - => guaranteed page utilization
 - => easy insertion/split algorithms.
 - (only deal with Minimum Bounding Rectangles - **MBRs**)






R-tree

- A multi-way external memory tree
- Index nodes and data (leaf) nodes
- All leaf nodes appear on the same level
- Every node contains between m and M entries
- The root node has at least 2 entries (children)

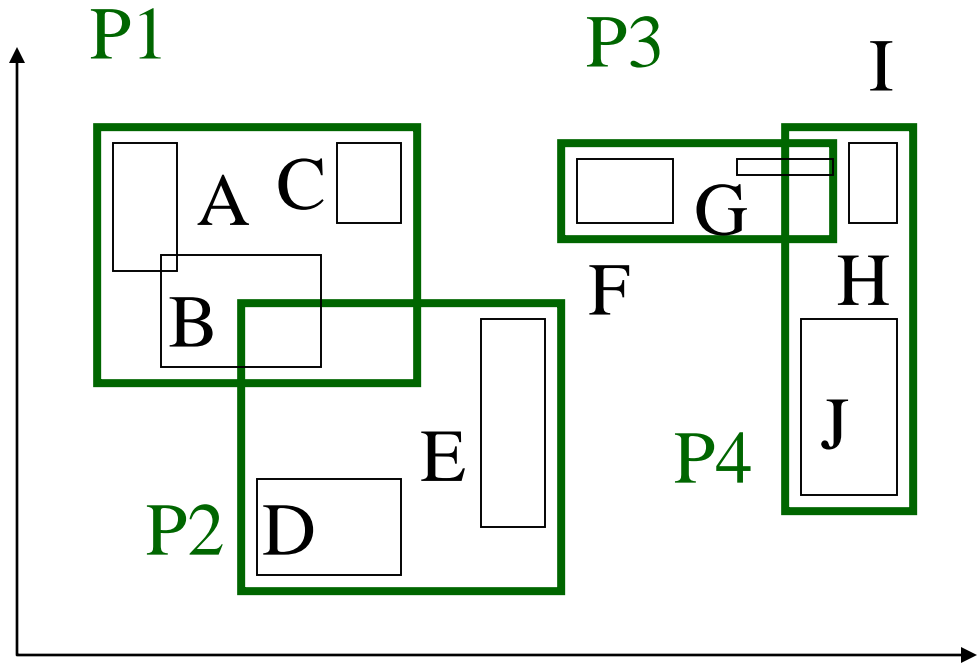
Example

- eg., w/ fanout 4: group nearby rectangles to parent MBRs; each group -> disk page



Example

■ $F=4$



A	B	C	
---	---	---	--


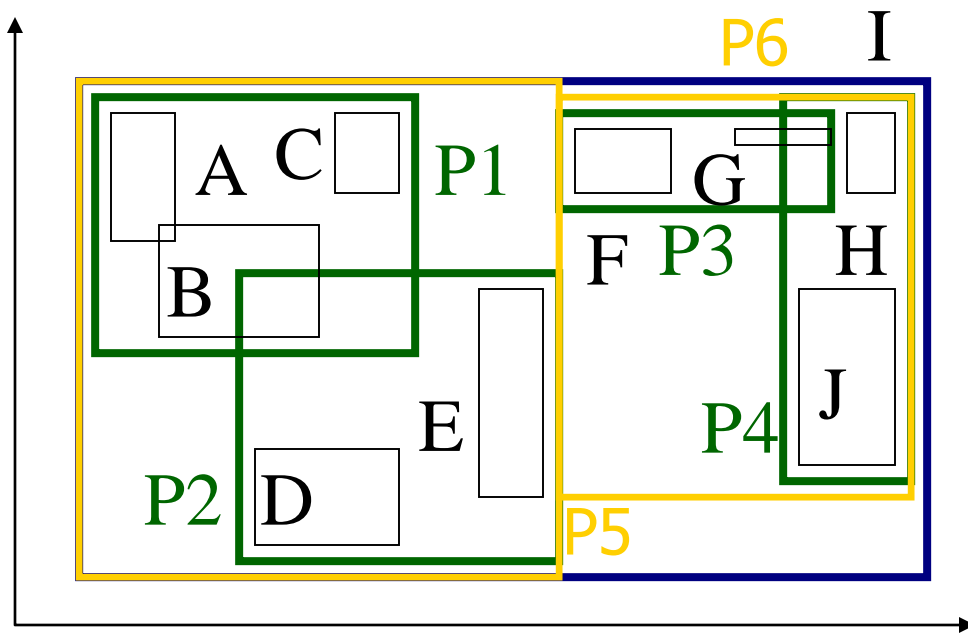
H	I	J	
---	---	---	--

D	E		
---	---	--	--

F	G		
---	---	--	--


Example

- $F=4 \mid m=2, M=4$




R-trees - format of nodes

- $\{(MBR; \text{obj_ptr})\}$ for leaf nodes


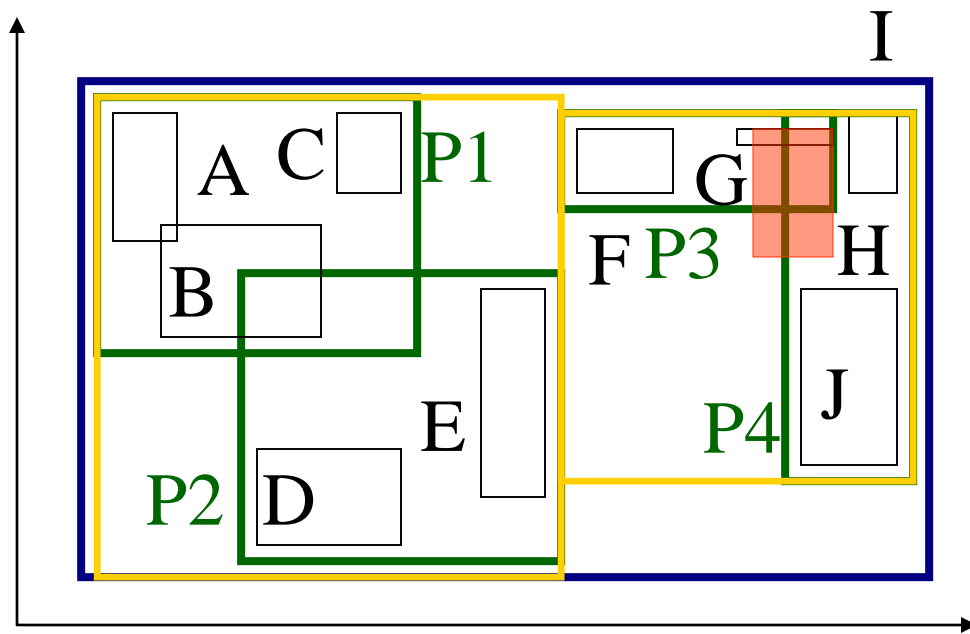


R-trees - format of nodes



- $\{(MBR; \text{node_ptr})\}$ for non-leaf nodes



R-trees: Search



R-trees: Search







R-trees:Search

- Main points:
 - every parent node completely covers its ‘children’
 - nodes in the same level may overlap!
 - a child MBR may be covered by more than one parent - it is stored under ONLY ONE of them. (ie., no need for dup. elim.)
 - a point query may follow multiple branches.
 - everything works for **any(?)** dimensionality


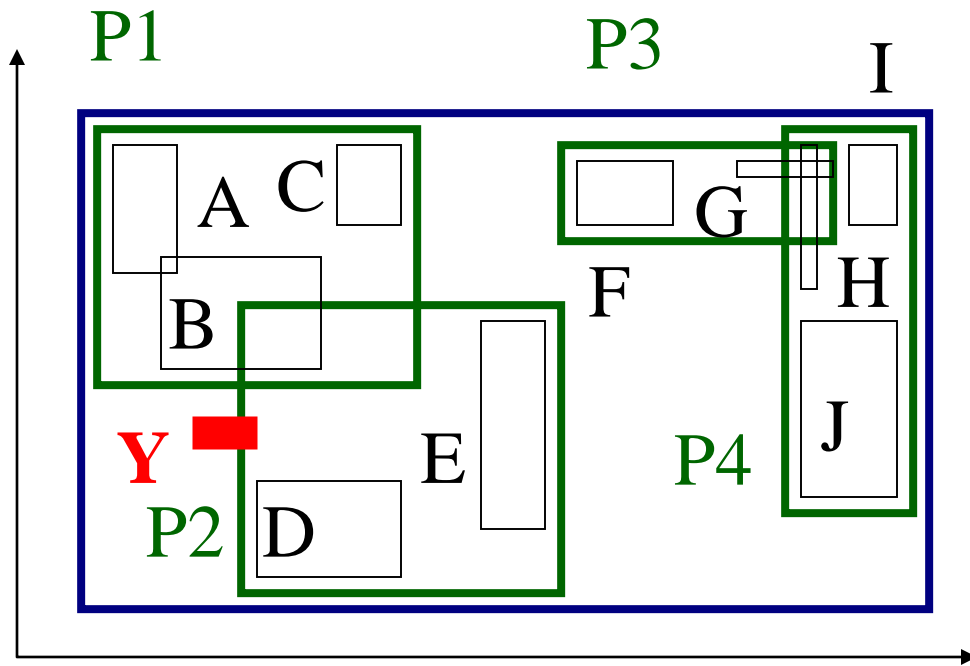
R-trees: Insertion

Insert X





R-trees: Insertion

Insert Y



R-trees: Insertion

- Extend the parent MBR







R-trees: Insertion

- How to find the next node to insert a new object Y?
 - Using ChooseLeaf: Find the entry that needs the least enlargement to include Y. Resolve ties using the area (smallest)
- Other methods (later)


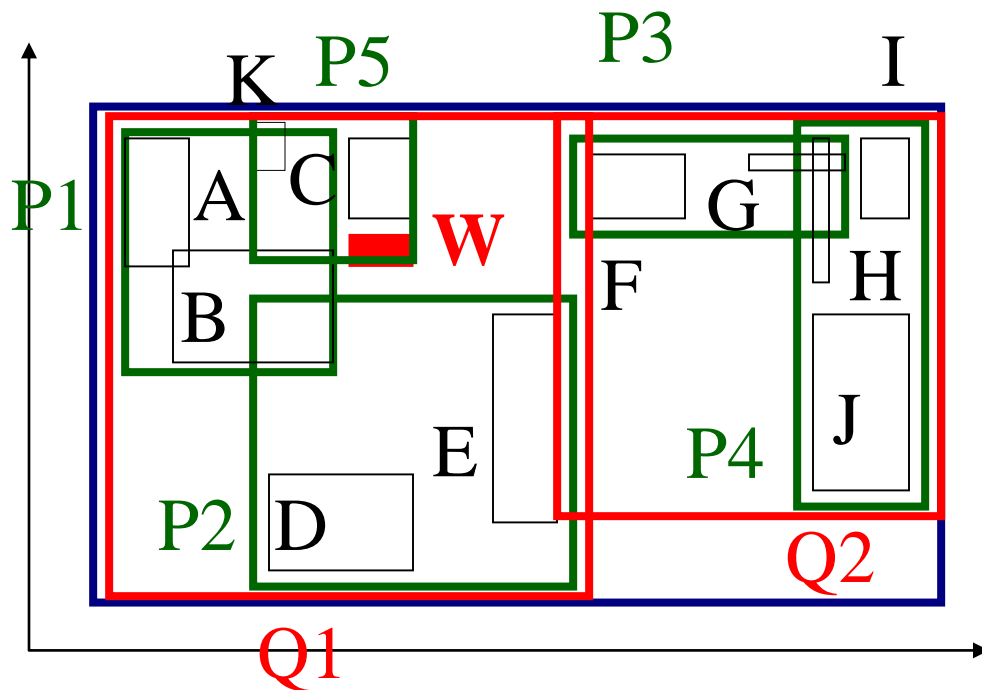
R-trees: Insertion

- If node is full then Split : ex. Insert w




R-trees: Insertion

- If node is full then Split : ex. Insert w



R-trees: Split

- Split node P1: partition the MBRs into two groups.




- (A1: plane sweep, until 50% of rectangles)
- A2: 'linear' split
- A3: quadratic split
- A4: exponential split:
 2^{M-1} choices




R-trees: Split

pick two rectangles as ‘seeds’ for group 1 and group 2;

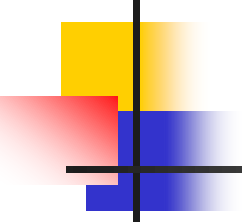



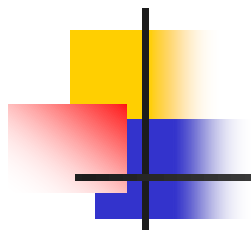
R-trees: Split

- pick two rectangles as 'seeds' for group 1 and group 2;
- assign each rectangle 'R' to the 'closest' 'group':
- 'closest': the smallest increase in area

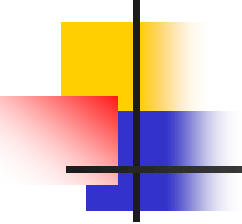


R-trees: Linear Split

- 
- How to pick Seeds:
 - Find the rects with the highest low and lowest high sides in each dimension
 - Normalize the separations by dividing by the width of all the rects in the corresponding dim
 - Choose the pair with the greatest normalized separation
 - PickNext: pick one of the remaining rects and add it to the closest group



R-trees: Quadratic Split

- 
- How to pick Seeds:
 - For each pair $E1$ and $E2$, calculate the rectangle $J = \text{MBR}(E1, E2)$ and $d = J - E1 - E2$. Choose the pair with the largest d
 - PickNext:
 - For each remaining rect calculate the area increase to include it in group $d1$ and $d2$
 - Choose rect with highest difference: $|d1 - d2|$
 - Assign to closest group



R-trees: Insertion

- Use the **ChooseLeaf** to find the leaf node to insert an entry E
- If leaf node is full, then **Split**, otherwise insert there
 - Propagate the split upwards, if necessary
- Adjust parent nodes



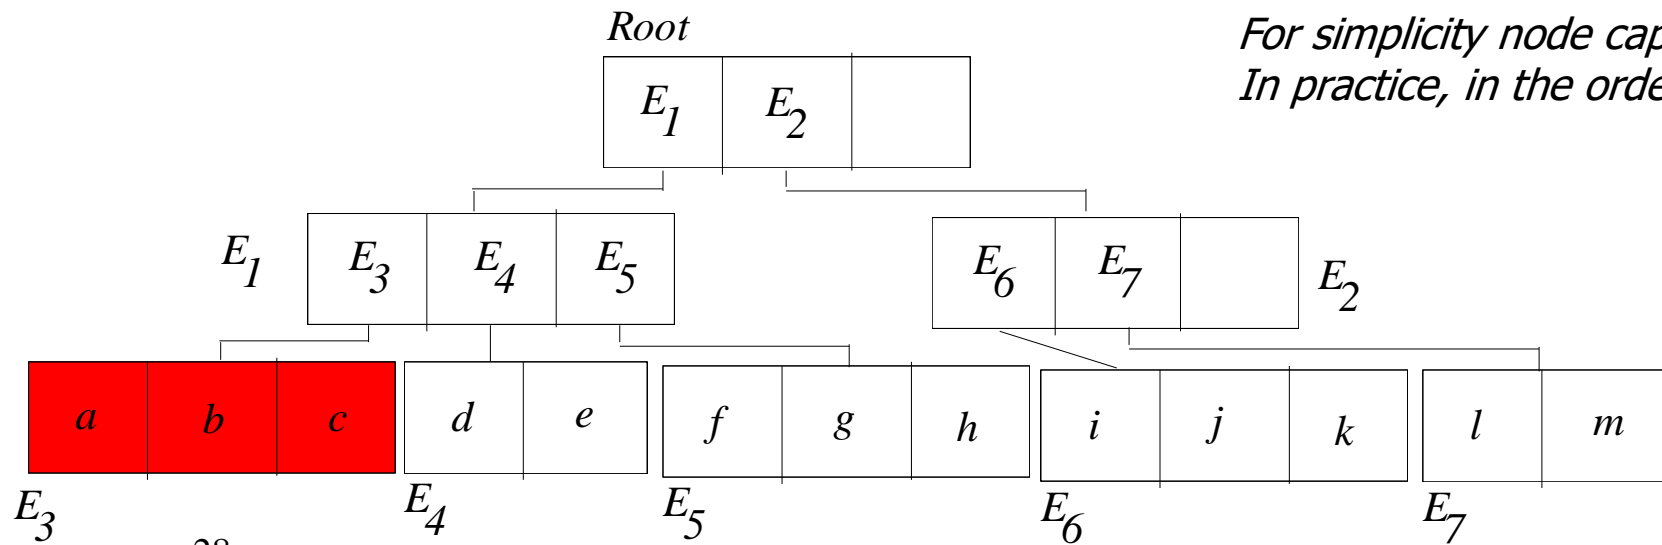

R-Trees: Deletion

- Find the leaf node that contains the entry E
- Remove E from this node
- If underflow:
 - Eliminate the node by removing the node entries and the parent entry
 - Reinsert the orphaned (other entries) into the tree using **Insert**
- Other method (**later**)




R-trees: Variations

- R+-tree: DO not allow overlapping, so split the objects (similar to z-values)
- R*-tree: change the insertion, deletion algorithms (minimize not only area but also perimeter, forced re-insertion)




*For simplicity node capacity = 3
In practice, in the order of 100*


R-Tree, Leaf Nodes




R-Tree – Intermediate Nodes



R-tree, Range Query



Range Query





R-tree

- The original R-tree tries to minimize the area of each enclosing rectangle in the index nodes.
- Is there any other property that can be optimized?

R*-tree → Yes!



R*-tree

- Optimization Criteria:
 - (O1) Area covered by an index MBR
 - (O2) Overlap between directory MBRs
 - (O3) Margin of a directory rectangle
 - (O4) Storage utilization
- Sometimes it is impossible to optimize all the above criteria at the same time!



R*-tree

- ChooseLeaf:

- If next node is not a leaf node, choose the node using the following criteria:
 - Least overlap enlargement
 - Least area enlargement
 - Smaller area
- Else
 - Least area enlargement
 - Smaller area

R*-tree



SplitNode

- Choose the axis to split
- Choose the two groups along the chosen axis
- ChooseSplitAxis
 - Along each axis, sort rectangles and break them into two groups ($M-2m+2$ possible ways where one group contains at least m rectangles). Compute the sum S of all margin-values (perimeters) of each pair of groups. Choose the one that minimizes S
- ChooseSplitIndex
 - Along the chosen axis, choose the grouping that gives the minimum overlap-value



R*-tree

- Forced Reinsert:
 - defer splits, by forced-reinsert, i.e.: instead of splitting, temporarily delete some entries, shrink overflowing MBR, and re-insert those entries
- Which ones to re-insert?
 - The ones that are further way from the center
- How many? A: 30%

