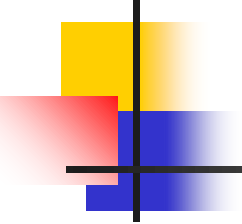




Temporal Databases

Outline

- 
-
- Spatial Databases
 - Indexing, Query processing
 - ■ Temporal Databases
 - Spatio-temporal
 -



Temporal DBs – Motivation

- Conventional databases represent the state of an enterprise at a single moment of time
- Many applications need information about the past
 - Financial (payroll)
 - Medical (patient history)
 - Government
- Temporal DBs: a system that manages time varying data

Comparison



- Conventional DBs:
 - Evolve through transactions from one state to the next
 - Changes are viewed as modifications to the state
 - No information about the past
 - Snapshot of the enterprise
- Temporal DBs:
 - Maintain historical information
 - Changes are viewed as additions to the information stored in the database
 - Incorporate notion of time in the system
 - Efficient access to past states

Temporal Databases



- Temporal Data Models: extension of relational model by adding temporal attributes to each relation
- Temporal Query Languages: TQUEL, SQL3
- Temporal Indexing Methods and Query Processing





Taxonomy of time

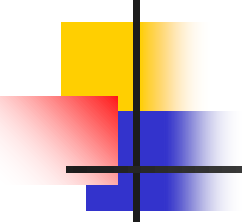
- Transaction time databases
 - Transaction time is the time when a fact is stored in the database
- Valid time databases:
 - Valid time is the time that a fact becomes effective in reality
- Bi-temporal databases:
 - Support both notions of time

Example



- Sales example: data about sales are stored at the end of the day
- Transaction time is different than valid time
- Valid time can refer to the future also!
 - Credit card: 03/13-04/16

Transaction Time DBs

- 
- Time evolves discretely, usually is associated with the transaction number:

$T1 \rightarrow T2 \rightarrow T3 \rightarrow T4 \dots$

- A record R is extended with an interval [t.start, t.end).
When we insert an object at t1 the temporal attributes are updated $\rightarrow [t1, \text{now})$
- Updates can be made only to the current state!
 - Past cannot be changed
 - “Rollback” characteristics


Transaction Time DBs

- Deletion is logical (never physical deletions!)
 - When an object is deleted at t_2 , its temporal attribute changes from $[t_1, \text{now}) \rightarrow [t_1, t_2)$ (lifetime)
 - Object is “alive” from insertion to deletion time, ex. t_1 to t_2 . If “now” then the object is still alive

eid	salary	start	end
10	20K	9/93	10/94
20	50K	4/94	*
33	30K	5/94	6/95
10	50K	1/95	*



Transaction Time DBs



Database evolves through insertions and deletions



Transaction Time DBs

- Requirements for index methods:
 - Store past logical states
 - Support addition/deletion/modification changes on the objects of the current state
 - Efficiently access and query any database state



Transaction Time DBs

- Queries:
 - Timestamp (timeslice) queries: ex. “Give me all employees at 05/94”
 - Range-timeslice: “Find all employees with id between 100 and 200 that worked in the company on 05/94”
 - Interval (period) queries: “Find all employees with id in [100,200] from 05/14 to 06/16”



Valid Time DBs


- Time evolves continuously
- Each object is a line segment representing its time span (eg. Credit card valid time)
- Support full operations on interval data:
 - Deletion at any time
 - Insertion at any time
 - Value change (modification) at any time (no ordering)

Valid Time DBs



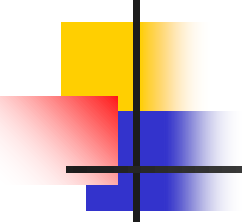
- Deletion is physical:
 - No way to know about the previous states of intervals
- The notion of “future”, “present” and “past” is relative to a certain timestamp t

Valid Time DBs



The reality “best know now !”

Valid Time DBs


- 
- Requirements for an Index method:
 - Store the latest collection of interval-objects
 - Support add/del/mod changes to this collection
 - Efficiently query the intervals in the collection
 - Timestamp query
 - Interval (period) query



Bitemporal DBs

- A transaction-time Database, but each record is an interval (plus the other attributes of the record)
- Keeping the evolution of a dynamic collection of interval-objects
- At each timestamp, it is a valid time database

Bitemporal DBs

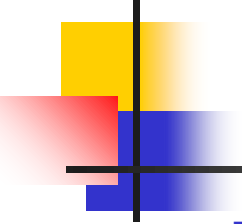




Bitemporal DBs

- Requirements for access methods:
 - Store past/logical states of collections of objects
 - Support add/del/mod of interval objects of the current logical state
 - Efficient query answering

Temporal Indexing

- 
- Straight-forward approaches:
 - B+-tree and R-tree
 - Problems?
 - Transaction time:
 - Snapshot Index, TSB-tree, **MVB-tree**, MVAS
 - Valid time:
 - Interval structures: Segment tree, even R-tree
 - Bitemporal:
 - Bitemporal R-tree



Temporal Indexing

- Lower bound on answering timeslice and range-timeslice queries:
 - Space $O(n/B)$, search $O(\log_B n + s/B)$
- n : number of changes, s : answer size, B page capacity

Range-timeslice: “Find all employees with id between 100 and 200 that worked in the company on 05/94”



Transaction Time Environment

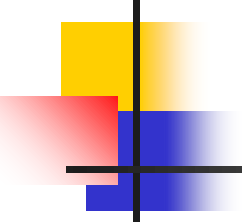
- Assume that when an event occurs in the real world it is inserted in the DB
- A timestamp is associated with each operation
- Transaction timestamps are monotonically increasing
- Previous transactions cannot be changed → we cannot change the past



Example

- Time evolving set of objects: employees of a company
- Time is discrete and described by a succession of non-negative integers: 1,2,3, ...
- Each time instant changes may happen, i.e., addition, deletion or modification
- We assume only insertion & deletion : modifications can be represented by a deletion and an insertion

Records

- 
- Each object is associated with:
 1. An oid (key, time invariant, eid)
 2. Attributes that can change (salary)
 3. A lifespan interval [t.start, t.end)
 - An object is alive from the time it inserted in the set, until it was deleted
 - At insertion time deletion is unknown
 - Deletions are logical: we change the now variable to the current time,
 $[t1, \text{now}) \rightarrow [t1, t2)$




Evolving set

- The state $S(t)$ of the evolving set at t is defined as: “the collection of alive objects at t ”
- The number of changes n represents a good metric for the minimal space needed to store the evolution

Evolving sets

- A new change updates the current state $S(t)$ to create a new state





Transaction-time Queries

- Pure-timeslice (time-point)
- Range-timeslice
- Pure-exact match



Snapshot Index

- Snapshot Index, is a method that answers efficiently pure-timeslice queries
- Based on a main memory method that solves the problem in $O(a + \log_2 n)$,
 $O(n)$ space
- External memory: $O(a/B + \log_B n)$


MM solution



- Copy approach: $O(a + \log n)$ but $O(n^2)$ space
- Log approach: $O(n)$ space but $O(n)$ query time
- We should combine the fast query time with the small space (and update)

Assumptions

- Assumptions (for clarity)
 - At each time instant there exist exactly one change
 - Each object is named by its creation time







Access Forest

- A double linked list L. Each new object is appended at the right end of L
- A deleted object is removed from L and becomes the next child of its left sibling in L
- Each object stores a pointer to its parent object. Also a parent object points to its first and last children
- So, each node has the following pointers:
parent, prev, next, Pcs, Pce

AF example





Additional structures

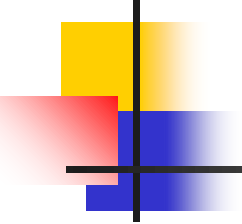
- A hashing scheme that keeps pointers to the positions of the alive elements in L
- An array A that stores the time changes. For **each time change instant**, it keeps the pointer to the current last object in L

Properties of AF

- In each tree of the AF the start times are sorted in preorder fashion
- The lifetime of an object includes the lifetimes of its descendants
- The intervals of two consecutive children under the same parent may have the following orderings:

$$s_i < e_i < s_{i+1} < e_{i+1} \text{ or } s_i < s_{i+1} < e_i < e_{i+1}$$

Searching

- 
- Find all objects alive at t_q
 - Use A to find the starting object in the access forest L ($O(\log n)$)
 - Traverse the access forest and report all alive objects at t_q $O(a)$ using the properties



Searching in AF

- Given query time q :
 - Use table A to find the time of the last object in L at time q . Say node Y .
 - Starting from Y go up (if it has a parent) recursively
 - For each node in the path from Y to the current node in list L (or Y itself if it is right now in L), visit the left sibling. If it is alive at q report it and recursively:
 - Visit the rightmost child in its subtree
 - Visit its left sibling
 - Stop the recursion overtime you find a node that is not alive at q

Disk based Solution



- Keep changes in pages as it was a Log
- Use hashing scheme to find objects by name (update $O(1)$)
- *Acceptor* : the current page that receives objects




Definitions

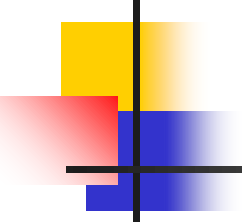
- A page is useful for the following time instants:
 - I-useful: while this page was the acceptor block
 - II-useful: for all time instants for which it contains at least uB “alive” records
- u is the usefulness parameter

Meta-evolution

- From the actual evolution of objects, now we have the evolution of pages! *meta-evolution*
- The “lifetime” of a page is its usefulness



Searching

- 
- Find all alive objects at $t_q \rightarrow$ Find all useful pages at t_q
 - The search can be done in $O(a/B + \log_B n)$



Copying procedure

- To maintain the answer in few pages we need clustering: controlled copying
- If a page has less than uB alive objects, we artificially delete the remaining alive objects and copy them to the acceptor block



Optimal Solution

- We can prove that the SI is optimal for pure-timeslice queries:
- $O(n)$ space, $O(a/B + \log_B n)$ query and $O(1)$ update (expected, using hashing)