



Spatio-Temporal Databases



Outline

- Spatial Databases
- Temporal Databases
- ■ Spatio-temporal Databases
- Multimedia Databases
-



Introduction

- Spatio-temporal Databases: manage spatial data whose geometry changes over time
- Geometry: position and/or extent

Examples:

- Global change data: climate or land cover changes
- Transportation: cars, airplanes
- Animated movies/video DBs



ST DBs


- A special Temporal Database
 - All the features of temporal database
 - Attributes can be spatial also
- Extension of Spatial Databases
 - Objects change instead of being static
 - At any timestamp it is a conventional Spatial Database
- New Database type




Requirements

- Efficient Representation of Space **and** Time
- Data Models
- Query Languages
- Query processing and Indexing
- GUI for spatio-temporal datasets

Spatio-temporal Objects



(a)
a moving point



(b)
a moving and shrinking region



ST Queries

- *Selection* Queries: “find all objects contained in a given area Q at a given time t ”
- *NN* queries: “find which object became the closest to a given point s during time interval T ,”
- *Aggregate* queries: “find how many objects passed through area Q during time interval T ,” or, “find the fastest object that will pass through area Q in the next 5 minutes from now”



ST Queries

- *join* queries: “given two spatiotemporal relations $R1$ and $R2$, find pairs of objects whose extents intersected during the time interval T ,” or “find pairs of planes that will come closer than 1 mile in the next 5 minutes”
- *similarity* queries: “find objects that moved similarly to the movement of a given object o over an interval T ”



ST Data Types

- Moving Points
 - Extent does not matter
 - Each object is modeled as a point (moving vehicles in a GIS based transportation system)
- Moving regions
 - Extent matters!
 - Each object is represented by an MBR, the MBR can change as the object move (airplanes, storm,...)




SP Data Types

- Different Type of changes:
 - Changes are applied discretely
 - Urban planning: appearance or dis-appearance of buildings
 - Changes are applied continuously
 - Moving objects (eg. Vehicles)

Trajectories

- Moving objects create trajectories
- Usually we can sample the positions of the objects at periodic time intervals Δt
- Linear Interpolation: easy and usually accurate enough
- Trajectory: a sequence of 2 or 3-dim locations






Temporal Environment

- Transaction or Valid time: (usually we assume transaction time)
- Two types of environments:
 - Predicting the future (or current) positions: Each object has a velocity vector. The DB can predict the location at any time $t \geq t_{\text{now}}$ assuming linear movement. Queries refer to the future
 - Storing the history. Queries refer to the past states of the spatial database


The Historical Environment

- Spatio-temporal Evolution



Indexing using R-trees

- Assume that time is another dimension, use a 3D R-tree
- Store the objects as their 3D MBR. How to compute that?





Problems of 3D R-tree

- How to store “now”? Use a large value...
- Common ending problem
- Long lived objects will have very long MBRs, difficult to cluster
- Extensive overlap and empty space → bad query performance for specific queries
- Also, works only for discrete changes




PPR-tree

- Better idea, partially persistent R-tree
- Two approaches: Multiversion and overlapping
- Multi-version: use the idea of the MVBT applied to R-tree.


Indexing Moving Objects

The problem of indexing any type of moving objects can be reduced to indexing **discrete rectangles**.


Discrete rectangles



Continuous points




Continuous rectangles



Historical R-trees (HR-trees)

An R-tree is maintained for each timestamp in history.

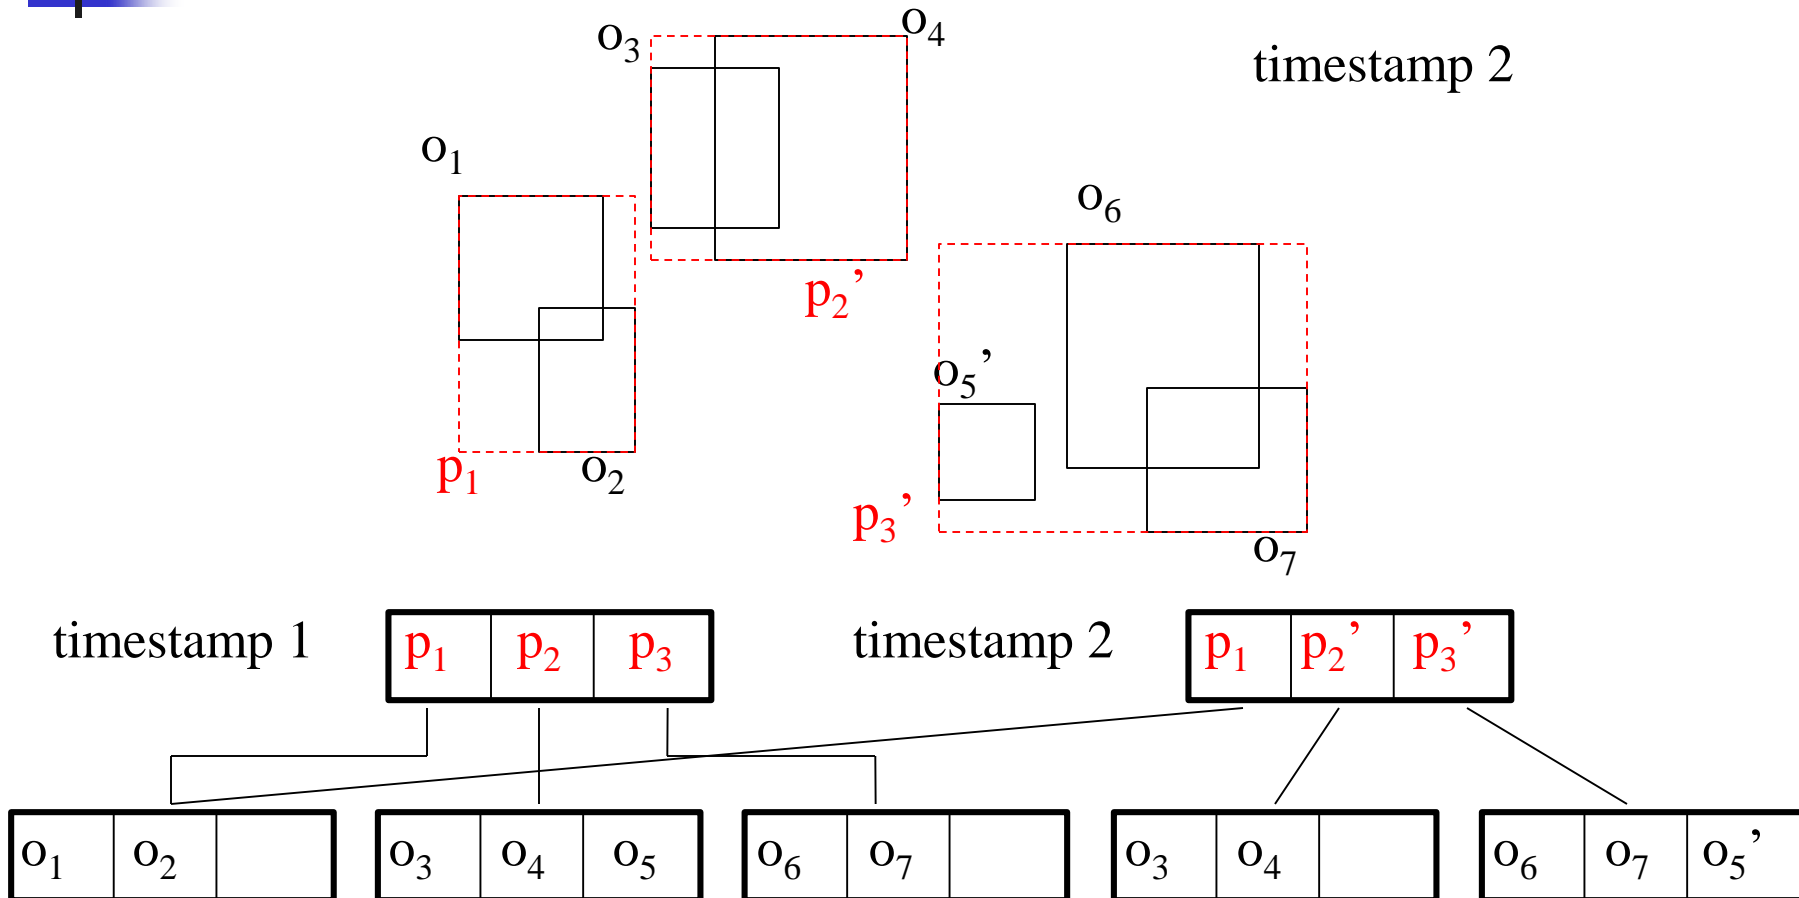
Trees at consecutive timestamps may share branches to save space.



Historical R-trees

An R-tree is maintained for each timestamp in history.

Trees at consecutive timestamps may share branches to save space.





HR-trees: Pros and Cons

- HR-trees answer timestamp queries very efficiently.
 - A timestamp query degenerates into a spatial window query handled by the corresponding R-tree at the query timestamp.
- Not quite efficient:
 - Expensive space consumption.
 - ➡ A node needs to be duplicated even when only one object moves.
 - Interval query processing is inefficient.
 - ➡ Although redundancy (from duplication) is necessary to maintain good timestamp query performance, it is excessive in HR-trees.



MVR-tree

- Consider a 2D R-tree that evolves over time
- Use the Multiversion B-tree approach to store the evolution of the tree...
- Need to consider some “spatial” issues:
 - No unique siblings, split methods, copies due to time split
- To insert a new object, compute a bounding box that encloses the object at all time instants, insert this bb as MBR



MVR-Tree

- An update or a query at some time instant t searches only among the spatial objects that are alive at t
- Space is linear to the number of updates, the problem of “now” is avoided
- Very efficient for snapshot or small interval queries



An improvement: The MV3R-tree

[Tao & Papadias 01]

- The MV3R-tree consists of a multi-version R-tree (MVR-tree) and an auxiliary 3D R-tree built on the leaves of the MVR-tree.
- The MVR-tree is optimized from the original multi-version framework, taking into account spatial properties.

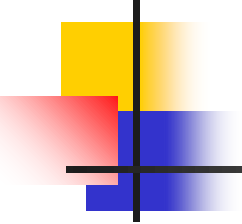


Redundancy in MVR-trees

- As with HR-trees, multi-version structures also aim at optimizing timestamp queries.
 - To achieve this goal, the original multi-version framework still leads to considerable redundancy, though much less than HR-trees.
- Excessive redundancy is harmful.
 - Space consumption is increased.
 - Compromise interval query performance, as multiple copies of the same object need to be retrieved.

Optimizing MVR-trees: Reducing Redundancy

- There are some heuristics to reduce data redundancy in MVR-trees in order to lower the space consumption and improve interval query performance.
 - Insertion
 1. General Key Split
 2. Insert in node after reinserting one of the entries
 3. Insert in another node
 4. Version split
- The resulting trees contain much less redundancy
 - Lower tree sizes and accelerate interval queries.
 - Timestamp queries are only slightly compromised.



Interval Query with Multi-Version Structures


- Need to search several logical trees responsible for the queried timestamps.
- Since multiple parent entries may point to the same child node, it is imperative to avoid duplicate accesses to the child.
 1. If a node is re-visited, the entire subtree rooted at the node needs to be re-visited.

Solution? Keep a hash table, other solutions exist also.



Building a 3D R-tree on the Leaves of the MVR-tree

- Build a 3D R-tree on the leaf nodes of the MVR-tree
- The size of the 3D R-tree is much smaller than a complete 3D R-tree as the number of leaf nodes is significantly lower than the number of actual objects.
- Since the 3D R-tree is not an acyclic graph but a tree, we do not have duplicate visit problem.
- Long interval queries are processed with the auxiliary 3D R-trees.



MVR-tree

leaf nodes

3D Rtree



What about moving objects?

- Problem: the MBR representation creates large **empty space**
- Use artificial deletes, approximate the object using many small MBRs
- But then, the space is increased
- Use an algorithm to distribute a small number of splits to the objects that need them most



What about moving objects?

- If objects move with linear functions of time:
Minimize total volume by splitting in
equidistant points
- Given K splits you can decide the best splits in
 $O(K \log N)$ time.

Reference:

[Tao & Papadias 01]:MV3R-Tree: A Spatio-Temporal Access Method for
Timestamp and Interval Queries. [VLDB 2001](#): 431-440