

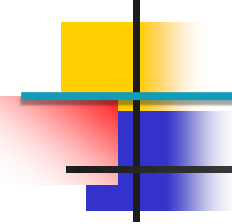


# LSI, SVD and Data Management

Based on the Slides from CS276: Information  
Retrieval and Web Search

Christopher Manning and Pandu Nayak

# Latent Semantic Indexing

- 
- Term-document matrices are very large
  - But the number of topics that people talk about is small (in some sense)
    - Clothes, movies, politics, ...
  - Can we represent the term-document space by a lower dimensional latent space?



---

---

# Linear Algebra Background

# Eigenvalues & Eigenvectors

- **Eigenvectors** (for a square  $m \times m$  matrix  $S$ )

$$S\mathbf{v} = \lambda\mathbf{v}$$

(right) eigenvector  $\mathbf{v} \in \mathbb{R}^m \neq \mathbf{0}$

eigenvalue  $\lambda \in \mathbb{R}$

*Example*

$$\begin{pmatrix} 6 & -2 \\ 4 & 0 \end{pmatrix} \begin{pmatrix} 1 \\ 2 \end{pmatrix} = \begin{pmatrix} 2 \\ 4 \end{pmatrix} = 2 \begin{pmatrix} 1 \\ 2 \end{pmatrix}$$

- How many eigenvalues are there at most?

$$S\mathbf{v} = \lambda\mathbf{v} \iff (S - \lambda\mathbf{I})\mathbf{v} = \mathbf{0}$$

only has a non-zero solution if  $|S - \lambda\mathbf{I}| = 0$

This is a  $m$ th order equation in  $\lambda$  which can have **at most  $m$  distinct solutions** (roots of the characteristic polynomial) - can be complex even though  $S$  is real.

# Matrix-vector multiplication

$S = \begin{pmatrix} 30 & 0 & 0 \\ 0 & 20 & 0 \\ 0 & 0 & 1 \end{pmatrix}$  has eigenvalues  $\lambda_1=30$ ,  $\lambda_2=20$ ,  $\lambda_3=1$  with corresponding eigenvectors

$$\vec{x}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}, \vec{x}_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \text{ and } \vec{x}_3 = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

On each eigenvector,  $S$  acts as a multiple of the identity matrix: but as a different multiple on each.

Any vector (say  $v = \begin{pmatrix} 2 \\ 4 \\ 6 \end{pmatrix}$ ) can be viewed as a combination of the eigenvectors:  $v = 2\vec{x}_1 + 4\vec{x}_2 + 6\vec{x}_3$

# Matrix-vector multiplication

- Thus a matrix-vector multiplication such as  $Sv$  ( $S$  matrix,  $v$  a vector) can be rewritten in terms of the eigenvalues/vectors:

$$\begin{aligned} S\vec{v} &= S(2\vec{x}_1 + 4\vec{x}_2 + 6\vec{x}_3) \\ &= 2S\vec{x}_1 + 4S\vec{x}_2 + 6S\vec{x}_3 \\ &= 2\lambda_1\vec{x}_1 + 4\lambda_2\vec{x}_2 + 6\lambda_3\vec{x}_3 \\ &= 60\vec{x}_1 + 80\vec{x}_2 + 6\vec{x}_3. \end{aligned}$$

- Even though  $v$  is an arbitrary vector, the action of  $S$  on  $v$  is determined by the eigenvalues/vectors.



# Matrix-vector multiplication

---

- Suggestion: the effect of “small” eigenvalues is small.
- If we ignored the smallest eigenvalue (1), then instead of

$$\begin{pmatrix} 60 \\ 80 \\ 6 \end{pmatrix} \quad \text{we would get} \quad \begin{pmatrix} 60 \\ 80 \\ 0 \end{pmatrix}$$

- These vectors are similar (in cosine similarity, etc.)

# Eigenvalues & Eigenvectors



For symmetric matrices, eigenvectors for distinct eigenvalues are **orthogonal**

$$Sv_{\{1,2\}} = \lambda_{\{1,2\}} v_{\{1,2\}} \text{ and } \lambda_1 \neq \lambda_2 \Rightarrow v_1^* v_2 = 0$$

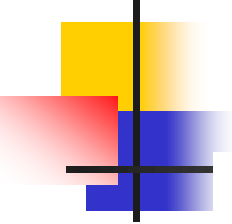
All eigenvalues of a real symmetric matrix are **real**.

All eigenvalues of a positive semidefinite symmetric matrix are **non-negative**

$$w^T S w \geq 0, \text{ for all } w$$



# Example



- Let  $S = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$  ← Real, symmetric.

- Then  $S - \lambda I = \begin{bmatrix} 2 - \lambda & 1 \\ 1 & 2 - \lambda \end{bmatrix} \Rightarrow$

$$|S - \lambda I| = (2 - \lambda)^2 - 1 = 0.$$

- The eigenvalues are 1 and 3 (nonnegative, real).
- The eigenvectors are orthogonal (and real):

$$\begin{pmatrix} 1 \\ -1 \end{pmatrix} \text{ and } \begin{pmatrix} 1 \\ 1 \end{pmatrix}.$$

# Eigen/diagonal Decomposition

- 
- Let  $S \in \mathbb{R}^{m \times m}$  be a symmetric **square** matrix with  **$m$  linearly independent eigenvectors** (a “non-defective” matrix)

- **Theorem:** Exists an **eigen decomposition**

$$S = U\Lambda U^{-1}$$

*diagonal*

Unique  
for  
distinct  
eigen-  
values

- (cf. matrix diagonalization theorem)
- Columns of **U** are the **eigenvectors** of **S**
- Diagonal elements of  $\Lambda$  are **eigenvalues** of **S**  
$$\Lambda = \text{diag}(\lambda_1, \dots, \lambda_m), \quad \lambda_i \geq \lambda_{i+1}$$



# Diagonal decomposition - example

---

Recall  $S = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$  with eigenvalues 3 and 1

The eigenvectors  $\begin{pmatrix} 1 \\ 1 \end{pmatrix}$  and  $\begin{pmatrix} 1 \\ -1 \end{pmatrix}$  define:  $U = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$

Inverting, we have  $U^{-1} = \begin{bmatrix} 1/2 & 1/2 \\ 1/2 & -1/2 \end{bmatrix}$

← Recall  
 $UU^{-1} = I.$

$$\text{Then, } \mathbf{S} = \mathbf{U} \mathbf{\Lambda} \mathbf{U}^{-1} = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} 3 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1/2 & 1/2 \\ 1/2 & -1/2 \end{bmatrix}$$

# Example continued

Let's divide  $\mathbf{U}$  (and multiply  $\mathbf{U}^{-1}$ ) by  $\sqrt{2}$

$$\text{Then, } \mathbf{S} = \underbrace{\begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix}}_{\mathbf{U}} \underbrace{\begin{bmatrix} 3 & 0 \\ 0 & 1 \end{bmatrix}}_{\neq} \underbrace{\begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix}}_{(\mathbf{U}^{-1} = \mathbf{U}^T)}$$

Why? Stay tuned ...

# Symmetric Eigen Decomposition

- If  $S \in \mathbb{R}^{m \times m}$  is a **symmetric** matrix:
- **Theorem**: There exists a (unique) **eigen decomposition**  $S = Q \Lambda Q^T$
- where **Q** is **orthogonal**:
  - $Q^{-1} = Q^T$
  - Columns of **Q** are normalized eigenvectors
  - Columns are orthogonal.
  - (everything is real)

# Singular Value Decomposition

For an  $M \times N$  matrix  $\mathbf{A}$  of rank  $r$  there exists a factorization (Singular Value Decomposition = **SVD**) as follows:

$$\mathbf{A} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$$

$M \times M$     $M \times N$     $V$  is  $N \times N$

(Not proven here.)

# Singular Value Decomposition

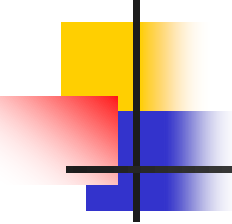

$$A = U \Sigma V^T$$

Diagram showing the dimensions of the matrices in the SVD equation:

- $U$  is  $M \times M$
- $\Sigma$  is  $M \times N$
- $V$  is  $N \times N$

- $AA^T = Q \Lambda Q^T$
- $AA^T = (U \Sigma V^T)(U \Sigma V^T)^T = (U \Sigma V^T)(V \Sigma U^T) = U \Sigma^2 U^T$

The columns of **U** are orthogonal eigenvectors of **AA<sup>T</sup>**.

The columns of **V** are orthogonal eigenvectors of **A<sup>T</sup>A**.

Eigenvalues  $\lambda_1 \dots \lambda_r$  of **AA<sup>T</sup>** are the eigenvalues of **A<sup>T</sup>A**.

$$\sigma_i = \sqrt{\lambda_i}$$
$$\Sigma = \text{diag}(\sigma_1, \dots, \sigma_r)$$

← Singular values

# Singular Value Decomposition


- Illustration of SVD dimensions and sparseness

$$\underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & \\ & \bullet & & \\ & & \bullet & \\ & & & \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_{V^T}$$

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & & \\ & \bullet & & \\ & & \bullet & \\ & & & \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{V^T}$$



# SVD example



Let  $A = \begin{bmatrix} 1 & -1 \\ 0 & 1 \\ 1 & 0 \end{bmatrix}$

Thus  $M=3$ ,  $N=2$ . Its SVD is

$$\begin{bmatrix} 0 & 2/\sqrt{6} & 1/\sqrt{3} \\ 1/\sqrt{2} & -1/\sqrt{6} & 1/\sqrt{3} \\ 1/\sqrt{2} & 1/\sqrt{6} & -1/\sqrt{3} \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & \sqrt{3} \\ 0 & 0 \end{bmatrix} \begin{bmatrix} 1/\sqrt{2} & 1/\sqrt{2} \\ 1/\sqrt{2} & -1/\sqrt{2} \end{bmatrix}$$

Typically, the singular values arranged in decreasing order.

# Low-rank Approximation

- SVD can be used to compute optimal **low-rank approximations**.
- Approximation problem: Find  $\mathbf{A}_k$  of rank  $k$  such that

$$\mathbf{A}_k = \min_{X: \text{rank}(X)=k} \|\mathbf{A} - X\|_F \quad \longleftarrow \text{Frobenius norm}$$

$$\|\mathbf{A}\|_F \equiv \sqrt{\sum_{i=1}^m \sum_{j=1}^n |a_{ij}|^2}.$$

$\mathbf{A}_k$  and  $X$  are both  $m \times n$  matrices.

Typically, want  $k \ll r$ .

# Low-rank Approximation

## ■ Solution via SVD

$$A_k = U \operatorname{diag}(\sigma_1, \dots, \sigma_k, \underbrace{0, \dots, 0}_{r-k}) V^T$$

*set smallest  $r-k$   
singular values to zero*

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{A_k} = \underbrace{\begin{bmatrix} * & * & \boxed{\phantom{0}} \\ * & * & \boxed{\phantom{0}} \\ * & * & \boxed{\phantom{0}} \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & \boxed{\phantom{0}} \\ & \bullet & \boxed{\phantom{0}} \\ \boxed{\phantom{0}} & & \boxed{\phantom{0}} \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ \hline \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} \\ \hline \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} & \phantom{0} \end{bmatrix}}_{V^T}$$

$$A_k = \sum_{i=1}^k \sigma_i u_i v_i^T$$



*column notation: **sum**  
of rank 1 matrices*

# Reduced SVD

- If we retain only  $k$  singular values, and set the rest to 0, then we don't need the matrix parts in color
- Then  $\Sigma$  is  $k \times k$ ,  $U$  is  $M \times k$ ,  $V^T$  is  $k \times N$ , and  $A_k$  is  $M \times N$
- This is referred to as the reduced SVD
- It is the convenient (space-saving) and usual form for computational applications
- It's what Matlab gives you

The diagram illustrates the Reduced SVD decomposition of a matrix  $A$  into three components:  $U$ ,  $\Sigma$ , and  $V^T$ .

Matrix  $A$  is shown as a 5x5 grid of stars. A blue bracket underneath the first three columns is labeled  $k$ , indicating that only the first  $k$  columns are retained in the reduced SVD.

Matrix  $U$  is shown as a 5x3 grid of stars. The third column is highlighted in blue, corresponding to the  $k$  columns of  $A$ .

Matrix  $\Sigma$  is shown as a 3x3 grid. The top-left element is a black dot, the middle element is a black dot, and the bottom-right element is a blue dot. The rest of the matrix is yellow, indicating that the singular values are zero.

Matrix  $V^T$  is shown as a 3x5 grid of stars. The first three rows are highlighted in blue, corresponding to the  $k$  rows of  $A$ .

The equation is represented as:

$$\underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_A = \underbrace{\begin{bmatrix} * & * & * \\ * & * & * \\ * & * & * \end{bmatrix}}_U \underbrace{\begin{bmatrix} \bullet & & \\ & \bullet & \\ & & \bullet \end{bmatrix}}_{\Sigma} \underbrace{\begin{bmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{bmatrix}}_{V^T}$$



# Approximation error

---

- How good (bad) is this approximation?
- It's the best possible, measured by the Frobenius norm of the error:

$$\min_{X: \text{rank}(X)=k} \|A - X\|_F$$

where the  $\sigma_i$  are ordered such that  $\sigma_i \geq \sigma_{i+1}$ .

Suggests why Frobenius error drops as  $k$  increases.



# SVD Low-rank approximation

---

- Whereas the term-doc matrix  $A$  may have  $M=50000$ ,  $N=10$  million (and rank close to 50000)
- We can construct an approximation  $A_{100}$  with rank 100.
  - Of all rank 100 matrices, it would have the lowest Frobenius error.
- Great ... but why would we??
- Answer: Latent Semantic Indexing



---

# Latent Semantic Indexing via the SVD



# What it is

---

- From term-doc matrix  $A$ , we compute the approximation  $A_k$ .
- There is a row for each term and a column for each doc in  $A_k$
- Thus docs live in a space of  $k \ll r$  dimensions
  - These dimensions are not the original axes
- But why?





# Vector Space Model: Pros

- **Automatic** selection of index terms
- **Partial matching** of queries and documents (dealing with the case where no document contains all search terms)
- **Ranking** according to **similarity score** (dealing with large result sets)
- **Term weighting** schemes (improves retrieval performance)
- Various extensions
  - Document clustering
  - Relevance feedback (modifying query vector)
- Geometric foundation



# Problems with Lexical Semantics

---

- Ambiguity and association in natural language
  - **Polysemy**: Words often have a **multitude of meanings** and different types of usage (more severe in very heterogeneous collections).
  - The vector space model is unable to discriminate between different meanings of the same word.

$$\text{sim}_{\text{true}}(d, q) < \cos(\angle(\vec{d}, \vec{q}))$$



# Latent Semantic Indexing (LSI)

---

- Perform a **low-rank approximation** of **document-term matrix** (typical rank **100–300**)
- General idea
  - Map documents (and terms) to a **low-dimensional** representation.
  - Design a mapping such that the low-dimensional space reflects **semantic associations** (latent semantic space).
  - Compute document similarity based on the **inner product** in this **latent semantic space**



# Goals of LSI

---

- LSI takes documents that are semantically similar (= talk about the same topics), but are not similar in the vector space (because they use different words) and re-represents them in a reduced vector space in which they have higher similarity.
- Similar terms map to similar location in low dimensional space
- Noise reduction by dimension reduction



## Example of $C = UV^T$ : The matrix $C$

$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	1	0	1	0	0	0
boat	0	1	0	0	0	0
ocean	1	1	0	0	0	0
wood	1	0	0	1	1	0
tree	0	0	0	1	0	1

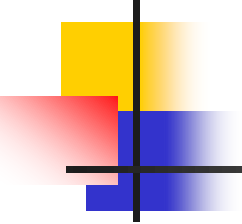
This is a typical term-document matrix. Actually, we use a non-weighted (binary) matrix here to simplify the example.



## Example of $C = UV^T$ : The matrix $U$

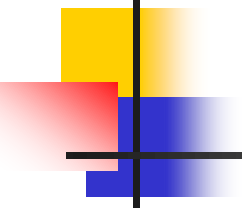
$U$	1	2	3	4	5
ship	-0.44	-0.30	0.57	0.58	0.25
boat	-0.13	-0.33	-0.59	0.00	0.73
ocean	-0.48	-0.51	-0.37	0.00	-0.61
wood	-0.70	0.35	0.15	-0.58	0.16
tree	-0.26	0.65	-0.41	0.58	-0.09

Example of  $C = U \Sigma V^T$  : The matrix  $\Sigma$



---

$\Sigma$	1	2	3	4	5
1	2.16	0.00	0.00	0.00	0.00
2	0.00	1.59	0.00	0.00	0.00
3	0.00	0.00	1.28	0.00	0.00
4	0.00	0.00	0.00	1.00	0.00
5	0.00	0.00	0.00	0.00	0.39



Example of  $C = U\Sigma V^T$  : The matrix  $V^T$

$V^T$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
1	-0.75	-0.28	-0.20	-0.45	-0.33	-0.12
2	-0.29	-0.53	-0.19	0.63	0.22	0.41
3	0.28	-0.75	0.45	-0.20	0.12	-0.33
4	0.00	0.00	0.58	0.00	-0.58	0.58
5	-0.53	0.29	0.63	0.19	0.41	-0.22



# Example of $C = U \Sigma V^T$ : All four matrices

$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$	
ship	1	0	1	0	0	0	
boat	0	1	0	0	0	0	
ocean	1	1	0	0	0	0	=
wood	1	0	0	1	1	0	
tree	0	0	0	1	0	1	
$U$	1	2	3	4	5		
ship	-0.44	-0.30	0.57	0.58	0.25		
boat	-0.13	-0.33	-0.59	0.00	0.73		
ocean	-0.48	-0.51	-0.37	0.00	-0.61		×
wood	-0.70	0.35	0.15	-0.58	0.16		
tree	-0.26	0.65	-0.41	0.58	-0.09		
$\Sigma$	1	2	3	4	5		
1	2.16	0.00	0.00	0.00	0.00		
2	0.00	1.59	0.00	0.00	0.00		
3	0.00	0.00	1.28	0.00	0.00		×
4	0.00	0.00	0.00	1.00	0.00		
5	0.00	0.00	0.00	0.00	0.39		
$V^T$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$	
1	-0.75	-0.28	-0.20	-0.45	-0.33	-0.12	
2	-0.29	-0.53	-0.19	0.63	0.22	0.41	
3	0.28	-0.75	0.45	-0.20	0.12	-0.33	
4	0.00	0.00	0.58	0.00	-0.58	0.58	
5	-0.53	0.29	0.63	0.19	0.41	-0.22	

# LSI: Summary



---

- We've decomposed the term-document matrix  $C$  into a product of three matrices.
- The term matrix  $U$  – consists of one (row) vector for each term
- The document matrix  $V^T$  – consists of one (column) vector for each document
- The singular value matrix  $\Sigma$  – diagonal matrix with singular values, reflecting importance of each dimension
- Next: Why are we doing this?

# How we use the SVD in LSI

- Key property: Each singular value tells us how important its dimension is.
- By setting less important dimensions to zero, we keep the important information, but get rid of the “details”.
- These details may
  - be **noise** – in that case, reduced LSI is a better representation because it is less noisy
  - **make things dissimilar that should be similar** – again reduced LSI is a better representation because it represents similarity better.
- Analogy for “fewer details is better”
  - Image of a bright red flower
  - Image of a black and white flower
  - Omitting color makes it easier to see similarity

# Reducing the dimensionality to 2


$U$	1	2	3	4	5	
ship	-0.44	-0.30	0.00	0.00	0.00	
boat	-0.13	-0.33	0.00	0.00	0.00	
ocean	-0.48	-0.51	0.00	0.00	0.00	
wood	-0.70	0.35	0.00	0.00	0.00	
tree	-0.26	0.65	0.00	0.00	0.00	
$\Sigma_2$	1	2	3	4	5	
1	2.16	0.00	0.00	0.00	0.00	
2	0.00	1.59	0.00	0.00	0.00	
3	0.00	0.00	0.00	0.00	0.00	
4	0.00	0.00	0.00	0.00	0.00	
5	0.00	0.00	0.00	0.00	0.00	
$V^T$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
1	-0.75	-0.28	-0.20	-0.45	-0.33	-0.12
2	-0.29	-0.53	-0.19	0.63	0.22	0.41
3	0.00	0.00	0.00	0.00	0.00	0.00
4	0.00	0.00	0.00	0.00	0.00	0.00
5	0.00	0.00	0.00	0.00	0.00	0.00

Actually, we only zero out singular values in  $\Sigma$ . This has the effect of setting the corresponding dimensions in  $U$  and  $V^T$  to zero when computing the product  $C = U \Sigma V^T$ .

# Reducing the dimensionality to 2

$C_2$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	0.85	0.52	0.28	0.13	0.21	-0.08
boat	0.36	0.36	0.16	-0.20	-0.02	-0.18
ocean	1.01	0.72	0.36	-0.04	0.16	-0.21
wood	0.97	0.12	0.20	1.03	0.62	0.41
tree	0.12	-0.39	-0.08	0.90	0.41	0.49
$U$	1	2	3	4	5	
ship	-0.44	-0.30	0.57	0.58	0.25	
boat	-0.13	-0.33	-0.59	0.00	0.73	
ocean	-0.48	-0.51	-0.37	0.00	-0.61	$\times$
wood	-0.70	0.35	0.15	-0.58	0.16	
tree	-0.26	0.65	-0.41	0.58	-0.09	
$\Sigma_2$	1	2	3	4	5	
1	2.16	0.00	0.00	0.00	0.00	
2	0.00	1.59	0.00	0.00	0.00	
3	0.00	0.00	0.00	0.00	0.00	$\times$
4	0.00	0.00	0.00	0.00	0.00	
5	0.00	0.00	0.00	0.00	0.00	
$V^T$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
1	-0.75	-0.28	-0.20	-0.45	-0.33	-0.12
2	-0.29	-0.53	-0.19	0.63	0.22	0.41
3	0.28	-0.75	0.45	-0.20	0.12	-0.33
4	0.00	0.00	0.58	0.00	-0.58	0.58
5	-0.53	0.29	0.63	0.19	0.41	-0.22

# Original matrix $C$ vs. reduced $C_2 = U \Sigma_2 V^T$



$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	1	0	1	0	0	0
boat	0	1	0	0	0	0
ocean	1	1	0	0	0	0
wood	1	0	0	1	1	0
tree	0	0	0	1	0	1
$C_2$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	0.85	0.52	0.28	0.13	0.21	-0.08
boat	0.36	0.36	0.16	-0.20	-0.02	-0.18
ocean	1.01	0.72	0.36	-0.04	0.16	-0.21
wood	0.97	0.12	0.20	1.03	0.62	0.41
tree	0.12	-0.39	-0.08	0.90	0.41	0.49

We can view  $C_2$  as a **two-dimensional** representation of the matrix. We have performed a **dimensionality reduction** to two dimensions.

# Why the reduced matrix is “better”

$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	1	0	1	0	0	0
boat	0	1	0	0	0	0
ocean	1	1	0	0	0	0
wood	1	0	0	1	1	0
tree	0	0	0	1	0	1
$C_2$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	0.85	0.52	0.28	0.13	0.21	-0.08
boat	0.36	0.36	0.16	-0.20	-0.02	-0.18
ocean	1.01	0.72	0.36	-0.04	0.16	-0.21
wood	0.97	0.12	0.20	1.03	0.62	0.41
tree	0.12	-0.39	-0.08	0.90	0.41	0.49

- Similarity of  $d_2$  and  $d_3$  in the original space: 0.  
 Similarity of  $d_2$  and  $d_3$  in the reduced space:  
 $0.52 * 0.28 + 0.36 * 0.16 + 0.72 * 0.36 + 0.12 * 0.20 + -0.39 * -0.08 \approx 0.52$

# Why the reduced matrix is “better”

$C$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	1	0	1	0	0	0
boat	0	1	0	0	0	0
ocean	1	1	0	0	0	0
wood	1	0	0	1	1	0
tree	0	0	0	1	0	1
$C_2$	$d_1$	$d_2$	$d_3$	$d_4$	$d_5$	$d_6$
ship	0.85	0.52	0.28	0.13	0.21	-0.08
boat	0.36	0.36	0.16	-0.20	-0.02	-0.18
ocean	1.01	0.72	0.36	-0.04	0.16	-0.21
wood	0.97	0.12	0.20	1.03	0.62	0.41
tree	0.12	-0.39	-0.08	0.90	0.41	0.49

“boat” and “ship” are semantically similar.  
The “reduced” similarity measure reflects this.



# Why we use LSI in information retrieval

- LSI takes documents that are semantically similar (= talk about the same topics), . . .
- . . . but are not similar in the vector space (because they use different words) . . .
- . . . and re-represents them in a reduced vector space . . .
- . . . in which they have higher similarity.
- Thus, LSI addresses the problems of **synonymy** and **semantic relatedness**.
- Standard vector space: Synonyms contribute nothing to document similarity.
- Desired effect of LSI: Synonyms contribute strongly to document similarity.

# How LSI addresses synonymy and semantic relatedness

- The dimensionality reduction forces us to omit a lot of “detail”.
- We have to map different words (= different dimensions of the full space) to the same dimension in the reduced space.
- The “cost” of mapping synonyms to the same dimension is much less than the cost of collapsing unrelated words.
- SVD selects the “least costly” mapping (see below).
- Thus, it will map synonyms to the same dimension.
- But it will avoid doing that for unrelated words.

# Implementation

- Compute SVD of term-document matrix
- Reduce the space and compute reduced document representations
- Map the query into the reduced space  
 $\vec{q}_2^T = \Sigma_2^{-1} U_2^T \vec{q}^T$   
 $C_2 = U \Sigma_2 V^T \Rightarrow \Sigma_2^{-1} U^T C = V_2^T$
- This follows from:
- Compute similarity of  $q_2$  with all reduced documents in  $V_2$ .
- Output ranked list of documents as usual
- Exercise: What is the fundamental problem with this approach?



# Resources

---

- Chapter 18 of IIR
- Resources at <http://ifnlp.org/ir>
  - Original paper on latent semantic indexing by Deerwester et al.
  - Paper on probabilistic LSI by Thomas Hofmann
  - Word space: LSI for words