

Extensions of Map Reduce Pig and Hive

Need for High-Level Languages

- Hadoop is great for large-data processing!
 - But writing Java programs for everything is verbose and slow
 - Data scientists don't want to write Java
- Solution: develop higher-level data processing languages
 - Hive: HQL is like SQL
 - Pig: Pig Latin is a bit like Perl

Hive and Pig

- Hive: data warehousing application in Hadoop
 - Query language is HQL, variant of SQL
 - Tables stored on HDFS with different encodings
 - Developed by Facebook, now open source
- Pig: large-scale data processing system
 - Scripts are written in Pig Latin, a dataflow language
 - Programmer focuses on data transformations
 - Developed by Yahoo!, now open source
- Common idea:
 - Provide higher-level language to facilitate large-data processing
 - Higher-level language “compiles down” to Hadoop jobs





Apache Pig

Based on slides from Adam Shook

What Is Pig?

- Developed by Yahoo! and a top level Apache project
- Immediately makes data on a cluster available to non-Java programmers via Pig Latin – a dataflow language
- Interprets Pig Latin and generates MapReduce jobs that run on the cluster
- Enables easy data summarization, ad-hoc reporting and querying, and analysis of large volumes of data
- Pig interpreter runs on a client machine – no administrative overhead required

Why Pig?

- Map reduce very low level!
- Need to write a lot of code and it is not interactive.
- Pig/Latin is introduced to address this

Pig Terms

- All data in Pig one of four types:
 - An Atom is a simple data value - stored as a string but can be used as either a string or a number
 - A Tuple is a data record consisting of a sequence of "fields"
 - Each field is a piece of data of any type (atom, tuple or bag)
 - A Bag is a set of tuples (also referred to as a 'Relation')
 - The concept of a "kind of a" table
 - A Map is a map from keys that are string literals to values that can be any data type
 - The concept of a hash map

Pig Capabilities

- Support for
 - Filtering
 - Grouping
 - Joins
 - Aggregation
- Extensibility
 - Support for User Defined Functions (UDF's)
- Leverages the same massive parallelism as native MapReduce

Pig Basics

- Pig is a client application
 - No cluster software is required
- Interprets Pig Latin scripts to MapReduce jobs
 - Parses Pig Latin scripts
 - Performs optimization
 - Creates execution plan
- Submits MapReduce jobs to the cluster

Execution Modes

- Pig has two execution modes
 - Local Mode - all files are installed and run using your local host and file system
 - MapReduce Mode - all files are installed and run on a Hadoop cluster and HDFS installation
- Interactive
 - By using the Grunt shell by invoking Pig on the command line

```
$ pig
grunt>
```
- Batch
 - Run Pig in batch mode using Pig Scripts and the "pig" command

```
$ pig -f id.pig -p <param>=<value> ...
```

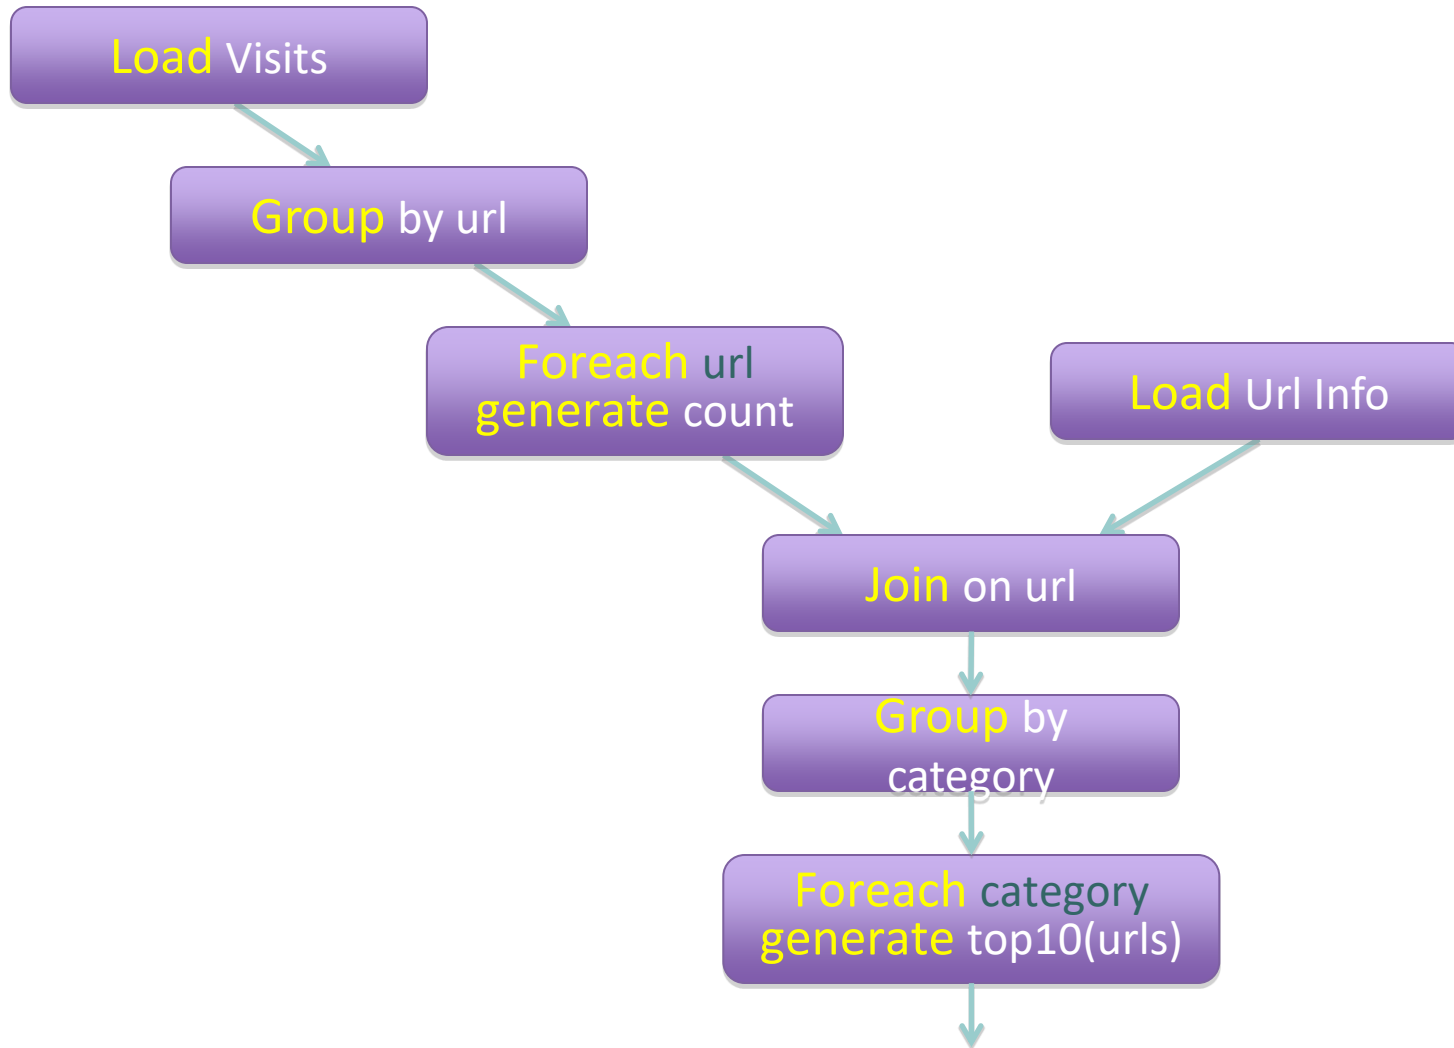
Pig Latin

- Pig Latin scripts are generally organized as follows
 - A LOAD statement reads data
 - A series of “transformation” statements process the data
 - A STORE statement writes the output to the filesystem
 - A DUMP statement displays output on the screen
- Logical vs. physical plans:
 - All statements are stored and validated as a logical plan
 - Once a STORE or DUMP statement is found the logical plan is executed

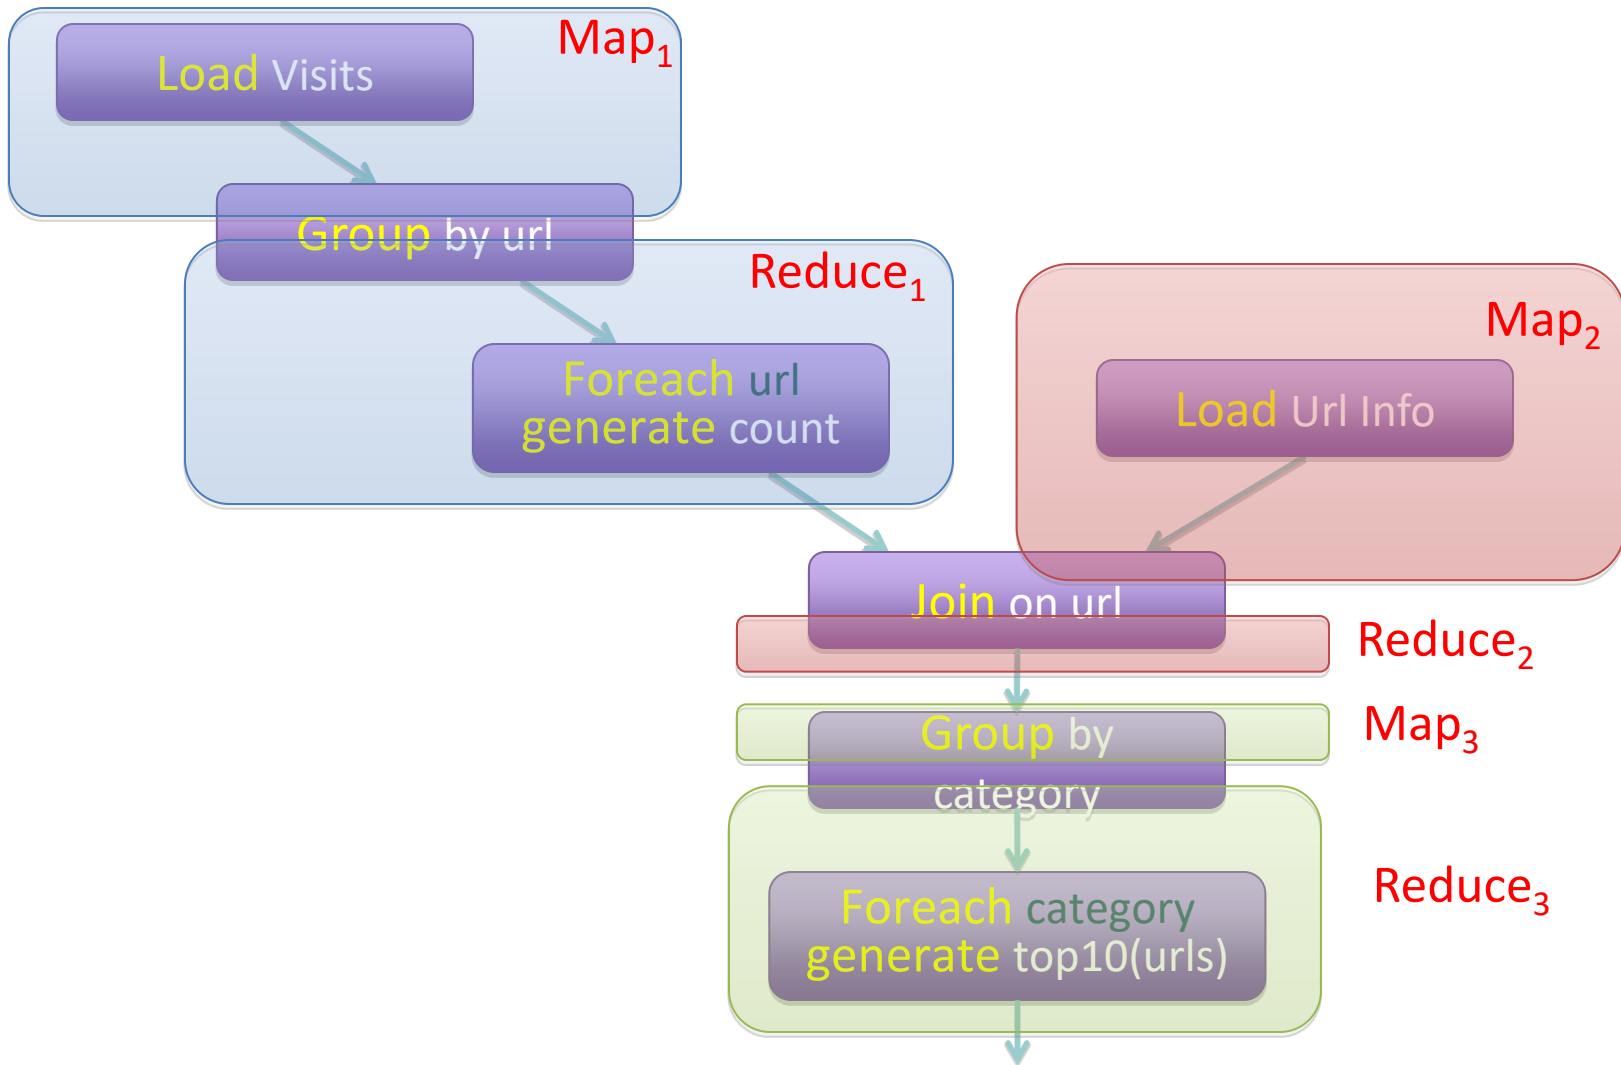
Pig Script

```
visits = load '/data/visits' as (user, url, time);  
gVisits = group visits by url;  
visitCounts = foreach gVisits generate url, count(visits);  
urlInfo = load '/data/urlInfo' as (url, category, pRank);  
visitCounts = join visitCounts by url, urlInfo by url;  
gCategories = group visitCounts by category;  
topUrls = foreach gCategories generate top(visitCounts,10);  
  
store topUrls into '/data/topUrls';
```

Pig Query Plan



Pig Script in Hadoop



Basic “grunt” Shell Commands

- Help is available

```
$ pig -h
```

- Pig supports HDFS commands

```
grunt> pwd
```

– put, get, cp, ls, mkdir, rm, mv, etc.

About Pig Scripts

- Pig Latin statements grouped together in a file
- Can be run from the command line or the shell
- Support parameter passing
- Comments are supported
 - Inline comments '--'
 - Block comments `/* */`

Simple Data Types

Type	Description
int	4-byte integer
long	8-byte integer
float	4-byte (single precision) floating point
double	8-byte (double precision) floating point
bytearray	Array of bytes; blob
chararray	String (“hello world”)
boolean	True/False (case insensitive)
datetime	A date and time
biginteger	Java BigInteger
bigdecimal	Java BigDecimal

Complex Data Types

Type	Description
Tuple	Ordered set of fields (a “row / record”)
Bag	Collection of tuples (a “resultset / table”)
Map	A set of key-value pairs Keys must be of type chararray

Pig Data Formats

- BinStorage
 - Loads and stores data in machine-readable (binary) format
- PigStorage
 - Loads and stores data as structured, field delimited text files
- TextLoader
 - Loads unstructured data in UTF-8 format
- PigDump
 - Stores data in UTF-8 format
- YourOwnFormat!
 - via UDFs

Loading Data Into Pig

- Loads data from an HDFS file

```
var = LOAD 'employees.txt';  
var = LOAD 'employees.txt' AS (id, name,  
    salary);  
var = LOAD 'employees.txt' using PigStorage()  
    AS (id, name, salary);
```
- Each LOAD statement defines a new bag
 - Each bag can have multiple elements (atoms)
 - Each element can be referenced by name or position (\$n)
- A bag is immutable
- A bag can be aliased and referenced later

Input And Output

- STORE

- Writes output to an HDFS file in a specified directory

```
grunt> STORE processed INTO 'processed_txt';
```

- Fails if directory exists
- Writes output files, part-[m|r]-xxxxx, to the directory

- PigStorage can be used to specify a field delimiter

- DUMP

- Write output to screen

```
grunt> DUMP processed;
```

Relational Operators

- FOREACH
 - Applies expressions to every record in a bag
- FILTER
 - Filters by expression
- GROUP
 - Collect records with the same key
- ORDER BY
 - Sorting
- DISTINCT
 - Removes duplicates

FOREACH ...GENERATE

- Use the FOREACH ...GENERATE operator to work with rows of data, call functions, etc.

- Basic syntax:

```
alias2 = FOREACH alias1 GENERATE expression;
```

- Example:

```
DUMP alias1;
```

```
(1,2,3) (4,2,1) (8,3,4) (4,3,3) (7,2,5) (8,4,3)
```

```
alias2 = FOREACH alias1 GENERATE col1, col2;
```

```
DUMP alias2;
```

```
(1,2) (4,2) (8,3) (4,3) (7,2) (8,4)
```

FILTER...BY

- Use the FILTER operator to restrict tuples or rows of data
- Basic syntax:

```
alias2 = FILTER alias1 BY expression;
```

- Example:

```
DUMP alias1;
```

```
(1,2,3) (4,2,1) (8,3,4) (4,3,3) (7,2,5) (8,4,3)
```

```
alias2 = FILTER alias1 BY (col1 == 8) OR (NOT  
    (col2+col3 > col1));
```

```
DUMP alias2;
```

```
(4,2,1) (8,3,4) (7,2,5) (8,4,3)
```


GROUP. . .ALL

- Use the GROUP...ALL operator to group data
 - Use GROUP when only one relation is involved
 - Use COGROUP with multiple relations are involved

- Basic syntax:

```
alias2 = GROUP alias1 ALL;
```

- Example:

```
DUMP alias1;
```

```
(John,18,4.0F) (Mary,19,3.8F) (Bill,20,3.9F)  
  (Joe,18,3.8F)
```

```
alias2 = GROUP alias1 BY col2;
```

```
DUMP alias2;
```

```
(18,{ (John,18,4.0F) , (Joe,18,3.8F) })  
(19,{ (Mary,19,3.8F) })  
(20,{ (Bill,20,3.9F) })
```

Pig: COGROUPing

```
A = LOAD 'myfile.txt' AS (f1: int, f2: int, f3: int);
```

A:

(1, 2, 3)

(4, 2, 1)

(8, 3, 4)

(4, 3, 3)

(7, 2, 5)

(8, 4, 3)

B:

(2, 4)

(8, 9)

(1, 3)

(2, 7)

(2, 9)

(4, 6)

(4, 9)

```
X = COGROUP A BY f1, B BY $0;
```

(1, {(1, 2, 3)}, {(1, 3)})

(2, {}, {(2, 4), (2, 7), (2, 9)})

(4, {(4, 2, 1), (4, 3, 3)}, {(4, 6), (4, 9)})

(7, {(7, 2, 5)}, {})

(8, {(8, 3, 4), (8, 4, 3)}, {(8, 9)})

ORDER...BY

- Use the ORDER...BY operator to sort a relation based on one or more fields

- Basic syntax:

```
alias = ORDER alias BY field_alias [ASC|DESC];
```

- Example:

```
DUMP alias1;
```

```
(1,2,3) (4,2,1) (8,3,4) (4,3,3) (7,2,5) (8,4,3)
```

```
alias2 = ORDER alias1 BY col3 DESC;
```

```
DUMP alias2;
```

```
(7,2,5) (8,3,4) (1,2,3) (4,3,3) (8,4,3) (4,2,1)
```

DISTINCT. . .

- Use the DISTINCT operator to remove duplicate tuples in a relation.
- Basic syntax:

```
alias2 = DISTINCT alias1;
```

- Example:

```
DUMP alias1;
```

```
(8,3,4) (1,2,3) (4,3,3) (4,3,3) (1,2,3)
```

```
alias2= DISTINCT alias1;
```

```
DUMP alias2;
```

```
(8,3,4) (1,2,3) (4,3,3)
```

Relational Operators

- FLATTEN
 - Used to un-nest tuples as well as bags
- INNER JOIN
 - Used to perform an inner join of two or more relations based on common field values
- OUTER JOIN
 - Used to perform left, right or full outer joins
- SPLIT
 - Used to partition the contents of a relation into two or more relations
- SAMPLE
 - Used to select a random data sample with the stated sample size

INNER JOIN. . .

- Use the JOIN operator to perform an inner, equi-join join of two or more relations based on common field values
- The JOIN operator always performs an inner join
- Inner joins ignore null keys
 - Filter null keys before the join
- JOIN and COGROUP operators perform similar functions
 - JOIN creates a flat set of output records
 - COGROUP creates a nested set of output records

INNER JOIN Example

```
DUMP Alias1;
```

```
(1,2,3)
```

```
(4,2,1)
```

```
(8,3,4)
```

```
(4,3,3)
```

```
(7,2,5)
```

```
(8,4,3)
```

```
DUMP Alias2;
```

```
(2,4)
```

```
(8,9)
```

```
(1,3)
```

```
(2,7)
```

```
(2,9)
```

```
(4,6)
```

```
(4,9)
```

```
Join Alias1 by Col1 to  
Alias2 by Col1
```

```
Alias3 = JOIN Alias1 BY  
Col1, Alias2 BY Col1;
```

```
Dump Alias3;
```

```
(1,2,3,1,3)
```

```
(4,2,1,4,6)
```

```
(4,3,3,4,6)
```

```
(4,2,1,4,9)
```

```
(4,3,3,4,9)
```

```
(8,3,4,8,9)
```

```
(8,4,3,8,9)
```

OUTER JOIN. . .

- Use the OUTER JOIN operator to perform left, right, or full outer joins
 - Pig Latin syntax closely adheres to the SQL standard
- The keyword OUTER is optional
 - keywords LEFT, RIGHT and FULL will imply left outer, right outer and full outer joins respectively
- Outer joins will only work provided the relations which need to produce nulls (in the case of non-matching keys) have schemas
- Outer joins will only work for two-way joins
 - To perform a multi-way outer join perform multiple two-way outer join statements

User-Defined Functions

- Natively written in Java, packaged as a jar file
 - Other languages include Jython, JavaScript, Ruby, Groovy, and Python
- Register the jar with the REGISTER statement
- Optionally, alias it with the DEFINE statement

```
REGISTER /src/myfunc.jar;  
A = LOAD 'students';  
B = FOREACH A GENERATE myfunc.MyEvalFunc($0);
```

DEFINE

- DEFINE can be used to work with UDFs and also streaming commands
 - Useful when dealing with complex input/output formats

```
/* read and write comma-delimited data */
DEFINE Y 'stream.pl' INPUT(stdin USING PigStreaming(', '))
    OUTPUT(stdout USING PigStreaming(', '));
A = STREAM X THROUGH Y;

/* Define UDFs to a more readable format */
DEFINE MAXNUM org.apache.pig.piggybank.evaluation.math.MAX;
A = LOAD 'student_data' AS (name:chararray, gpa1:float, gpa2:double);
B = FOREACH A GENERATE name, MAXNUM(gpa1, gpa2);
DUMP B;
```

Example Pig Script

```
-- Load the content of a file into a pig bag named 'input_lines'
input_lines = LOAD 'CHANGES.txt' AS (line:chararray);

-- Extract words from each line and put them into a pig bag named 'words'
words = FOREACH input_lines GENERATE FLATTEN(TOKENIZE(line)) AS word;

-- filter out any words that are just white spaces
filtered_words = FILTER words BY word MATCHES '\\w+';

-- create a group for each word
word_groups = GROUP filtered_words BY word;

-- count the entries in each group
word_count = FOREACH word_groups GENERATE COUNT(filtered_words) AS count, group AS word;

-- order the records by count
ordered_word_count = ORDER word_count BY count DESC;

-- Store the results ( executes the pig script )
STORE ordered_word_count INTO 'output';
```

PageRank in Pig

```
previous_pagerank = LOAD '$docs_in' USING PigStorage()  
  AS (url: chararray, pagerank: float,  
    links:{link: (url: chararray)});
```

```
outbound_pagerank = FOREACH previous_pagerank  
  GENERATE pagerank / COUNT(links) AS pagerank,  
  FLATTEN(links) AS to_url;
```

```
new_pagerank =  
  FOREACH ( COGROUP outbound_pagerank  
    BY to_url, previous_pagerank BY url INNER )  
  GENERATE group AS url,  
    (1 - $d) + $d * SUM(outbound_pagerank.pagerank) AS pagerank,  
    FLATTEN(previous_pagerank.links) AS links;
```

```
STORE new_pagerank INTO '$docs_out' USING PigStorage();
```

Oh, the iterative part too...

```
#!/usr/bin/python
from org.apache.pig.scripting import *
P = Pig.compile(""" Pig part goes here """)

params = { 'd': '0.5', 'docs_in': 'data/pagerank_data_simple' }

for i in range(10):
    out = "out/pagerank_data_" + str(i + 1)
    params["docs_out"] = out
    Pig.fs("rmr " + out)
    stats = P.bind(params).runSingle()
    if not stats.isSuccessful():
        raise 'failed'
    params["docs_in"] = out
```

References

- <http://pig.apache.org>



Apache Hive

Based on Slides by Adam Shook

What Is Hive?

- Developed by Facebook and a top-level Apache project
- A data warehousing infrastructure based on Hadoop
- Immediately makes data on a cluster available to non-Java programmers via SQL like queries
- Built on HiveQL (HQL), a SQL-like query language
- Interprets HiveQL and generates MapReduce jobs that run on the cluster
- Enables easy data summarization, ad-hoc reporting and querying, and analysis of large volumes of data

What Hive Is Not

- Hive, like Hadoop, is designed for batch processing of large datasets
- Not an OLTP or real-time system
- Latency and throughput are both high compared to a traditional RDBMS
 - Even when dealing with relatively small data (<100 MB)

Data Hierarchy

- Hive is organised hierarchically into:
 - Databases: namespaces that separate tables and other objects
 - Tables: homogeneous units of data with the same schema
 - Analogous to tables in an RDBMS
 - Partitions: determine how the data is stored
 - Allow efficient access to subsets of the data
 - Buckets/clusters
 - For sub-sampling within a partition
 - Join optimization

HiveQL

- HiveQL / HQL provides the basic SQL-like operations:
 - Select columns using SELECT
 - Filter rows using WHERE
 - JOIN between tables
 - Evaluate aggregates using GROUP BY
 - Store query results into another table
 - Download results to a local directory (i.e., export from HDFS)
 - Manage tables and queries with CREATE, DROP, and ALTER

Primitive Data Types

Type	Comments
TINYINT, SMALLINT, INT, BIGINT	1, 2, 4 and 8-byte integers
BOOLEAN	TRUE/FALSE
FLOAT, DOUBLE	Single and double precision real numbers
STRING	Character string
TIMESTAMP	Unix-epoch offset <i>or</i> datetime string
DECIMAL	Arbitrary-precision decimal
BINARY	Opaque; ignore these bytes

Complex Data Types

Type	Comments
STRUCT	A collection of elements If S is of type STRUCT {a INT, b INT}: S.a returns element a
MAP	Key-value tuple If M is a map from 'group' to GID: M['group'] returns value of GID
ARRAY	Indexed list If A is an array of elements ['a','b','c']: A[0] returns 'a'

HiveQL Limitations

- HQL only supports equi-joins, outer joins, left semi-joins
- Because it is only a shell for Map-Reduce, complex queries can be hard to optimise
- Missing large parts of full SQL specification:
 - HAVING clause in SELECT
 - Correlated sub-queries
 - Sub-queries outside FROM clauses
 - Updatable or materialized views
 - Stored procedures

Hive Metastore

- Stores Hive metadata
- Default metastore database uses Apache Derby
- Various configurations:
 - Embedded (in-process metastore, in-process database)
 - Mainly for unit tests
 - Local (in-process metastore, out-of-process database)
 - Each Hive client connects to the metastore directly
 - Remote (out-of-process metastore, out-of-process database)
 - Each Hive client connects to a metastore server, which connects to the metadata database itself

Hive Warehouse

- Hive tables are stored in the Hive “warehouse”
 - Default HDFS location: /user/hive/warehouse
- Tables are stored as sub-directories in the warehouse directory
- Partitions are subdirectories of tables
- External tables are supported in Hive
- The actual data is stored in flat files

Hive Schemas

- Hive is schema-on-read
 - Schema is only enforced when the data is read (at query time)
 - Allows greater flexibility: same data can be read using multiple schemas
- Contrast with an RDBMS, which is schema-on-write
 - Schema is enforced when the data is loaded
 - Speeds up queries at the expense of load times

Create Table Syntax

```
CREATE TABLE table_name  
    (col1 data_type,  
     col2 data_type,  
     col3 data_type,  
     col4 datatype )  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS format_type;
```

Simple Table

```
CREATE TABLE page_view
  (viewTime INT,
   userid BIGINT,
   page_url STRING,
   referrer_url STRING,
   ip STRING COMMENT 'IP Address of the User' )
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE;
```

More Complex Table

```
CREATE TABLE employees (  
    (name STRING,  
    salary FLOAT,  
    subordinates ARRAY<STRING>,  
    deductions MAP<STRING, FLOAT>,  
    address STRUCT<street:STRING,  
                    city:STRING,  
                    state:STRING,  
                    zip:INT>)  
  
    ROW FORMAT DELIMITED  
    FIELDS TERMINATED BY '\t'  
    STORED AS TEXTFILE;
```

External Table

```
CREATE EXTERNAL TABLE page_view_stg
  (viewTime INT,
   userid BIGINT,
   page_url STRING,
   referrer_url STRING,
   ip STRING COMMENT 'IP Address of the User')
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE
LOCATION '/user/staging/page_view';
```

More About Tables

- CREATE TABLE
 - LOAD: file moved into Hive's data warehouse directory
 - DROP: both metadata and data deleted
- CREATE EXTERNAL TABLE
 - LOAD: no files moved
 - DROP: only metadata deleted
 - Use this when sharing with other Hadoop applications, or when you want to use multiple schemas on the same data

Partitioning

- Can make some queries faster
- Divide data based on partition column
- Use PARTITION BY clause when creating table
- Use PARTITION clause when loading data
- SHOW PARTITIONS will show a table's partitions

Bucketing

- Can speed up queries that involve sampling the data
 - Sampling works without bucketing, but Hive has to scan the entire dataset
- Use **CLUSTERED BY** when creating table
 - For sorted buckets, add **SORTED BY**
- To query a sample of your data, use **TABLESAMPLE**

Browsing Tables And Partitions

Command	Comments
<code>SHOW TABLES;</code>	Show all the tables in the database
<code>SHOW TABLES 'page.*';</code>	Show tables matching the specification (uses regex syntax)
<code>SHOW PARTITIONS page_view;</code>	Show the partitions of the page_view table
<code>DESCRIBE page_view;</code>	List columns of the table
<code>DESCRIBE EXTENDED page_view;</code>	More information on columns (useful only for debugging)
<code>DESCRIBE page_view PARTITION (ds='2008-10-31');</code>	List information about a partition

Loading Data

- Use LOAD DATA to load data from a file or directory
 - Will read from HDFS unless LOCAL keyword is specified
 - Will append data unless OVERWRITE specified
 - PARTITION required if destination table is partitioned

```
LOAD DATA LOCAL INPATH '/tmp/pv_2008-06-8_us.txt'  
OVERWRITE INTO TABLE page_view  
PARTITION (date='2008-06-08', country='US')
```

Inserting Data

- Use INSERT to load data from a Hive query
 - Will append data unless OVERWRITE specified
 - PARTITION required if destination table is partitioned

```
FROM page_view_stg pvs
  INSERT OVERWRITE TABLE page_view
  PARTITION (dt='2008-06-08', country='US')
  SELECT pvs.viewTime, pvs.userid,
         pvs.page_url, pvs.referrer_url
  WHERE pvs.country = 'US';
```

Loading And Inserting Data: Summary

Use this	For this purpose
LOAD	Load data from a file or directory
INSERT	Load data from a query <ul style="list-style-type: none">• One partition at a time• Use multiple INSERTs to insert into multiple partitions in the one query
CREATE TABLE AS (CTAS)	Insert data while creating a table
Add/modify external file	Load new data into external table

Sample Select Clauses

- Select from a single table

```
SELECT *  
  FROM sales  
 WHERE amount > 10 AND  
        region = "US";
```

- Select from a partitioned table

```
SELECT page_views.*  
  FROM page_views  
 WHERE page_views.date >= '2008-03-01' AND  
        page_views.date <= '2008-03-31'
```

Relational Operators

- ALL and DISTINCT
 - Specify whether duplicate rows should be returned
 - ALL is the default (all matching rows are returned)
 - DISTINCT removes duplicate rows from the result set
- WHERE
 - Filters by expression
 - Does not support IN, EXISTS or sub-queries in the WHERE clause
- LIMIT
 - Indicates the number of rows to be returned

Relational Operators

- GROUP BY
 - Group data by column values
 - Select statement can only include columns included in the GROUP BY clause
- ORDER BY / SORT BY
 - ORDER BY performs total ordering
 - Slow, poor performance
 - SORT BY performs partial ordering
 - Sorts output from each reducer

- Hive looks similar to an SQL database
- Relational join on two tables:
 - Table of word counts from Shakespeare collection
 - Table of word counts from the bible

```
SELECT s.word, s.freq, k.freq FROM shakespeare s
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1
ORDER BY s.freq DESC LIMIT 10;
```

the	25848	62394
I	23031	8854
and	19671	38985
to	18038	13526
of	16700	34654
a	14170	8057
you	12702	2720
my	11297	4135
in	10797	12445
is	8882	6884

Hive: Behind the Scenes

```
SELECT s.word, s.freq, k.freq FROM shakespear s
JOIN bible k ON (s.word = k.word) WHERE s.freq >= 1 AND k.freq >= 1
ORDER BY s.freq DESC LIMIT 10;
```



(Abstract Syntax Tree)

```
(TOK_QUERY (TOK_FROM (TOK_JOIN (TOK_TABREF shakespear s) (TOK_TABREF bible k) (= (. (TOK_TABLE_OR_COL s) word) (.
(TOK_TABLE_OR_COL k) word)))) (TOK_INSERT (TOK_DESTINATION (TOK_DIR TOK_TMP_FILE)) (TOK_SELECT (TOK_SELEXPR (.
(TOK_TABLE_OR_COL s) word)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL s) freq)) (TOK_SELEXPR (. (TOK_TABLE_OR_COL k) freq))) (TOK_WHERE
(AND (>= (. (TOK_TABLE_OR_COL s) freq) 1) (>= (. (TOK_TABLE_OR_COL k) freq) 1))) (TOK_ORDERBY (TOK_TABSORTCOLNAMEDESC (.
(TOK_TABLE_OR_COL s) freq))) (TOK_LIMIT 10)))
```



(one or more of MapReduce jobs)

Hive: Behind the Scenes

STAGE DEPENDENCIES:

Stage-1 is a root stage

Stage-2 depends on stages: Stage-1

Stage-0 is a root stage

STAGE PLANS:

Stage: Stage-1

Map Reduce

Alias -> Map Operator Tree:

s

TableScan

alias: s

Filter Operator

predicate:

expr: (freq >= 1)

type: boolean

Reduce Output Operator

key expressions:

expr: word

type: string

sort order: +

Map-reduce partition columns:

expr: word

type: string

tag: 0

value expressions:

expr: freq

type: int

expr: word

type: string

k

TableScan

alias: k

Filter Operator

predicate:

expr: (freq >= 1)

type: boolean

Reduce Output Operator

key expressions:

expr: word

type: string

sort order: +

Map-reduce partition columns:

expr: word

type: string

tag: 1

value expressions:

expr: freq

type: int

Reduce Operator Tree:

Join Operator

condition map:

Inner Join 0 to 1

condition expressions:

0 {VALUE._col0} {VALUE._col1}

1 {VALUE._col0}

outputColumnNames: _col0, _col1, _col2

Filter Operator

predicate:

expr: ((_col0 >= 1) and (_col2 >= 1))

type: boolean

Select Operator

expressions:

expr: _col1

type: string

expr: _col0

type: int

expr: _col2

type: int

outputColumnNames: _col0, _col1, _col2

File Output Operator

compressed: false

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.SequenceFileInputFormat

output format: org.apache.hadoop.hive.ql.io.HiveSequenceFileOutputFormat

Stage: Stage-2

Map Reduce

Alias -> Map Operator Tree:

hdfs://localhost:8022/tmp/hive-training/364214370/10002

Reduce Output Operator

key expressions:

expr: _col1

type: int

sort order: -

tag: -1

value expressions:

expr: _col0

type: string

expr: _col1

type: int

expr: _col2

type: int

Reduce Operator Tree:

Extract

Limit

File Output Operator

compressed: false

GlobalTableId: 0

table:

input format: org.apache.hadoop.mapred.TextInputFormat

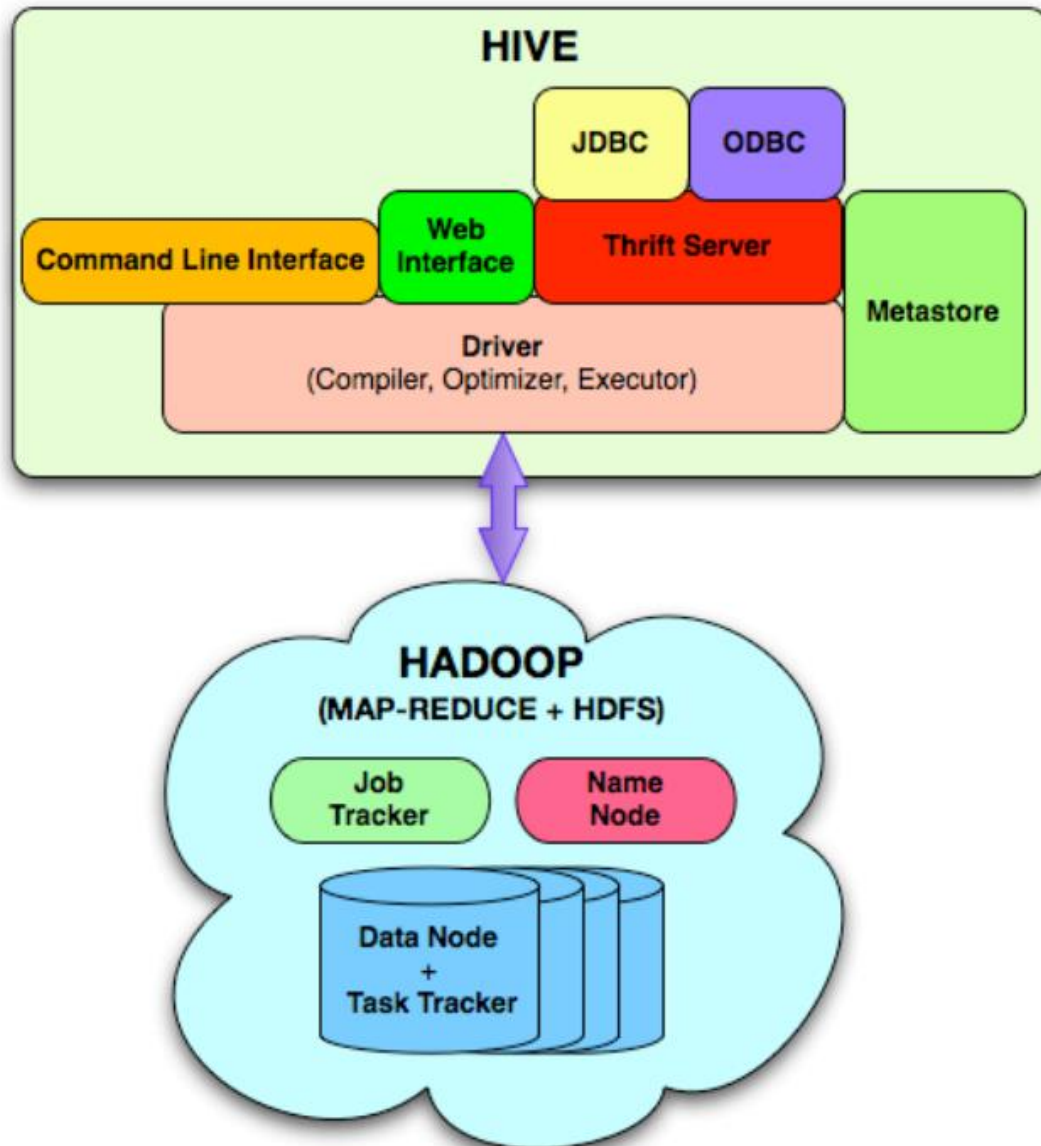
output format: org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat

Stage: Stage-0

Fetch Operator

limit: 10

Hive Architecture



References

- <http://hive.apache.org>